

Earth as a Field of Admissible Futures

Geoengineering, Biomass, and Civic Simulation

Flyxion

Independent Researcher

Draft — September 5, 2026

Abstract

Geoengineering discussion is conventionally organized around a single image: an apparatus acting on a single planetary variable, most often temperature, to move it toward a preferred value. This book argues that the image is inadequate on its own terms. No optimal global temperature exists to be targeted, since the systems a temperature might be optimized for, biomass, biodiversity, agriculture, coastal stability, ocean productivity, respond to warming differently and do not share a maximum; nor does maximizing total biomass serve as an adequate fallback, since biomass in the wrong place, concentration, or state can represent ecological damage rather than health regardless of its quantity. In place of a single scalar objective, this book develops a continuation-geometric standard, reachability, viability, recoverability, and path dependence, and argues that evaluating interventions against it requires replacing narrative coordination, discourse, persuasion, symbolic alignment, with shared, persistent simulation, in which competing proposals are explored as executable hypotheses before any commitment is made.

That method is applied across four intervention families: active polar albedo maintenance, where a nuclear-refrigeration proposal and a lower-intensity cloud-and-salt hybrid trade concentrated leverage against distributed reversibility; oceanic continuation engineering, where iceberg routing, kelp and Sargassum processing, and a thermopneumatic conversion reactor show mobility itself functioning as a resource, purchased at the cost of continuously renegotiated rather than fixed-site admissibility; constructed terrestrial productivity, where support architectures expanding a rainforest's own growth capacity are evaluated by a graded standard tracking their progress toward ecological autonomy rather than by biomass added; and planetary material circulation, where siting, routing, and full material and energy chain tracing expose tensions between interventions that no single intervention's own evaluation could surface. A closing part addresses authorization, evidence, and repair directly: who may commit to an intervention, why the population of proposals a society actually evaluates is a biased survivor sample rather than a fair one, why termination planning belongs at the point of initial evaluation rather than as an afterthought, and what a civilization actually retains from an intervention's failure.

The result is not a blueprint for which future Earth should have, but a standard, and a method for applying it, for judging which futures remain reachable after a given intervention's choices have been made.

Contents

Introduction: The Planet Has No Single Set Point	3
I Beyond the Planetary Thermostat	5
1 The Fiction of the Optimal Temperature	6
2 Biomass Is Not Biospheric Health	9
3 Climate Change as Commitment Mismatch	13
4 Continuation Geometry	16
II Seeing Before Acting	20
6 Simulation Before Commitment	21
7 Civilizations as Executable Hypotheses	25
III Active Polar Albedo Maintenance	31
8 Poles as Control Surfaces	32
9 The Refrigerator-Ice-Machine Proposal	35
10 Holding the Seams: Cloud-and-Salt Hybrid Albedo	38
IV Oceanic Continuation Engineering	42
11 Boundary Conditions in Motion	43
12 Icebergs as Mobile Boundary Conditions	45
13 Kelp Farms and the Sargassum Problem	48
14 The Caldera Reactor: Thermopneumatic Transformation of Maladmitted Biomass	51

V	Constructed Terrestrial Productivity	55
15	Substrate Before Growth	56
16	Rainforest Generators: Expanding the Growth State Space	58
17	Terra Preta Rain and the Marine-Terrestrial Loop	61
VI	Planetary Material Circulation	64
5	Placement Without Disposal	65
18	Orthodromic Infrastructure: Routing Between Reserves	70
19	Tracing the Full Material and Energy Chain	73
20	Coupled Intervention Portfolios	75
VII	Governance, Failure, and Repair	78
21	Record, Recognize, Authorize, Commit, Consequence	79
22	Evidence Before Recognition: What Survives Is Not What Was Best	81
23	Termination Without Termination Shock	84
24	Repair as Planetary Knowledge	87
	Conclusion: Reversible Supports for Plural Continuations	90

Introduction: The Planet Has No Single Set Point

Geoengineering is usually imagined through a single image: a machine, or a coordinated set of machines, acting on a single planetary variable to move it toward a preferred value. The variable is almost always temperature, whether the mechanism is stratospheric aerosol injection, solar radiation management, or a nuclear-powered installation of the kind Neal Stephenson imagines in *Termination Shock*. In that novel, a private actor builds an intervention large enough to alter the climate unilaterally, forcing every affected government, institution, and population to respond to a fact rather than deliberate over a proposal. The novel is not a source of evidence for anything argued here, but it dramatizes something true about the ordinary shape of the geoengineering imagination: one apparatus, one variable, one planet moved toward one target.

This book replaces that image with a different one. It treats geoengineering as a portfolio of interventions distributed across water, ice, clouds, ocean currents, biomass, nutrients, and infrastructure, each acting on a different part of the Earth system, each subject to its own constraints and failure modes, and none of them reducible to a single planetary thermostat. Chapter 1 has already shown why the thermostat model fails even on its own terms: there is no established optimal global temperature, because the systems that might be optimized (biomass, biodiversity, agricultural productivity, coastal stability, ocean circulation) respond to temperature differently and do not share a maximum. Chapter 2 has shown why the most obvious fallback objective, maximizing total biomass, fails in the same way, since biomass existing in the wrong place or beyond a receiving system's capacity to incorporate it can represent damage rather than health regardless of its quantity.

Given those two results, the central question this book asks is not “what temperature should Earth be” but “which climatic and ecological futures should remain reachable, for whom, under what risks, and with what capacity for correction.” This question does not collapse into a scalar, and it is not meant to. It is meant to keep visible a plurality of outcomes, populations, and ecosystems that a single target necessarily suppresses, and to hold every intervention this book considers, including the ones it proposes, to the standard of preserving that plurality rather than collapsing it toward one preferred future.

The book's remaining parts follow the shape of that question rather than the shape of a single mechanism. Part I, of which this introduction and the following four chapters are the opening movement, establishes the conceptual ground: the fiction of the optimal temperature, the distinction between biomass and biospheric health, climate change as a mismatch between committed systems and revised regularities, and the vocabulary of continuation geometry (reachability, viability, recoverability, and path dependence) that the rest of the book uses to evaluate specific proposals. Part II turns from concept to method, asking how competing interventions, models, and preferences can be explored and compared before any of them is committed to at planetary scale; it does this by adapting two ex-

isting frameworks, *Simulation as Civic Infrastructure* and *Civilizations as Hypotheses*, to the specific problem of planetary intervention, rather than developing a new theory of governance from scratch. Parts III through VI then work through specific interventions in turn: active polar albedo maintenance, oceanic continuation engineering, constructed terrestrial productivity, and planetary material circulation, each held to the standards Chapters 1 and 2 establish rather than judged by whether it lowers a single number. Part VII closes with governance, failure, and repair, asking not only who is authorized to alter a shared atmosphere but how interventions can be designed from the outset to be terminated, corrected, or reversed without producing a second crisis in the act of undoing the first.

Stephenson's Schmidt gets one thing right that this book preserves: once an intervention becomes materially credible, it forces positions that abstract deliberation could defer indefinitely. This book's proposals aim for the same forcing function without the same unilateralism, by using shared simulation, explicit governance structure, and termination planning to make positions visible and contestable before, rather than after, an intervention becomes irreversible.

Part I

Beyond the Planetary Thermostat

Chapter 1

The Fiction of the Optimal Temperature

Climate policy is conventionally organized around a single number: the increase in global mean surface temperature relative to a preindustrial baseline. Targets are expressed in fractions of a degree. Success or failure is reported as a threshold crossed or held. This framing has been indispensable for coordinating international attention, and it will remain indispensable as a diagnostic variable throughout this book. But a diagnostic variable is not the same thing as an engineering objective, and the temptation to treat it as one has quietly organized a great deal of geoengineering discussion around a question that has no defensible answer: what temperature should Earth be?

No such temperature has been established, and there is good reason to think none exists in the form the question presupposes. The difficulty is not that the answer is currently unknown but likely to be found with better models. The difficulty is that “optimal” requires an objective, and the objectives that might be optimized are neither singular nor aligned with one another.

Consider a partial list of the quantities a planetary temperature affects: total biomass, biodiversity, agricultural productivity, human heat tolerance, disease range, coastal land area, ocean primary production, ice sheet stability, freshwater availability, and the frequency of extreme weather. Each of these responds to temperature differently, and not always monotonically. A given warming increment may lengthen a growing season in one latitude while shortening it in another. It may increase net photosynthesis in a boreal forest while triggering dieback in a tropical one. It may improve conditions for one commercially important fish species while degrading them for another sharing the same water column. There is no reason to expect these responses to share a maximum, and considerable reason to expect they do not.

This is not a claim that all temperatures are equally good, or that climate change is not a serious problem. It is a claim about the shape of the problem. The question “is this level of warming dangerous” can be answered, often clearly, once a specific system and a specific form of danger are named: coastal infrastructure committed to a stable sea level, agricultural regions committed to a rainfall pattern, coral reefs committed to a narrow thermal tolerance. The question “what is the optimal global temperature” cannot be answered in the same way, because it asks for a single scalar to summarize outcomes that are irreducibly plural and frequently in tension.

A particularly instructive case is the relationship between temperature and biomass, because it is often assumed, implicitly, that warming is either straightforwardly bad for the biosphere or straightforwardly good for plant growth, and neither assumption survives close examination. Some of the most biologically productive regions in the present ocean are cold, not because coldness directly produces life, but because low temperatures coincide with circulation regimes that deliver nutrients to the sur-

face. Winds and ocean circulation bring cold, nutrient-rich water upward into sunlit layers where phytoplankton can use it. This upwelling, rather than coldness by itself, makes the Southern Ocean and the eastern boundary currents of the Pacific exceptionally productive marine environments. A warmer ocean can therefore hold more thermal energy while supporting less primary production in exactly the regions where upwelling currently sustains large food webs.

This does not establish that a colder Earth would be biologically superior to a warmer one. It establishes that the relationship between temperature and productivity runs through intermediate variables, chiefly wind-driven circulation and nutrient delivery, that do not scale simply with temperature in either direction. Warming in one region can open new nutrient pathways even as it closes others elsewhere. Both very warm and very cold planetary states have supported extensive marine productivity under the right circulation conditions, and neither state guarantees it.

The same caution applies to land. Extended growing seasons and elevated atmospheric carbon dioxide can increase certain crop yields in specific regions, holding water and nutrients constant, which they typically are not. Warming also raises evapotranspiration, shifts rainfall patterns, expands the range of some agricultural pests and pathogens, and pushes many staple crops toward or past their thermal tolerance during critical growth stages. Whether a given region experiences net agricultural benefit or net agricultural loss from a given warming increment depends on its starting climate, its crops, its water infrastructure, and the rate at which the change occurs relative to the capacity of farmers, seed varieties, and irrigation systems to adapt. There is no global answer, only a distribution of regional ones, and averaging that distribution would conceal precisely the differences an intervention must preserve.

If temperature does not determine a single correct value for biomass, biodiversity, or productivity, it follows that biomass itself cannot serve as a simpler proxy objective in its place. This might seem to invite a fallback position: perhaps the goal should simply be to maximize the total quantity of living matter the planet supports, since more biomass would seem to indicate a more vigorous biosphere. Chapter 2 develops in detail why this fallback fails, using two paired cases, the early closure of the Antarctic krill fishery and the record-breaking growth of the Great Atlantic Sargassum Belt, to show that biomass existing in the wrong place, at the wrong concentration, or beyond a receiving system's capacity to incorporate it can represent ecological damage rather than ecological health, regardless of the total tonnage involved. The short version of that argument, needed here only to close off the fallback, is that quantity and function are different measurements. A planet can gain biomass while losing the ecological relationships that made a smaller quantity of it valuable.

This has direct consequences for how climate risk should be framed. It is common, and not wrong, to note that a large share of visible climate damage concentrates on low-lying coastal populations facing sea-level rise. But treating coastal flooding as the central problem risks understating everything else at stake, and it also risks implying that a strategy calibrated to protect coastlines has thereby addressed climate change. A more general and more accurate framing is that climate risk arises wherever human or ecological systems have become committed, often over decades or centuries, to a set of climatic regularities that a changing climate is now revising. Coastal cities are one instance of this commitment, made visible by sea walls, property values, and evacuation maps. Agricultural regions committed to a rainfall pattern are another, less visible until a multi-year drought exposes it. Coral reefs committed to a narrow thermal tolerance are a third, revealed by bleaching events rather than infrastructure failure. Fisheries committed to a seasonal ice cycle, water utilities committed to

a snowpack-fed river, and insurance markets committed to historical loss tables are further instances of the same underlying structure. Chapter 3 develops this commitment-mismatch framing directly; it is introduced here only to show why coastal flooding, however serious, is a special case of a general problem rather than the problem itself.

None of this implies that a single global temperature target is useless as a coordinating device. Diagnostic variables earn their keep by being measurable, comparable across regions and years, and correlated with a wide range of underlying risks, and global mean temperature satisfies all three conditions well enough to remain central to monitoring and to international agreement. The error is not in measuring it. The error is in treating its reduction as identical to the achievement of planetary health, and in treating any proposed intervention primarily by whether it moves that one number in the preferred direction.

The chapters that follow adopt a different governing question. Rather than asking what temperature Earth should have, they ask which climatic and ecological trajectories should remain reachable, for which populations and ecosystems, under what risks, and with what capacity for correction if the intervention turns out to be mistaken. This question does not resolve into a single target, and it is not meant to. It is meant to keep visible the plurality of outcomes that a single scalar necessarily suppresses, and to hold every intervention considered in this book, including the ones proposed here, to the standard of preserving that plurality rather than collapsing it toward one preferred future.

Chapter 2

Biomass Is Not Biospheric Health

In 2025, the Antarctic krill fishery reached its annual 620,000-tonne trigger level early for the first time. It did so again in 2026. At nearly the same time, the tropical Atlantic was producing Sargassum on an extraordinary scale. The Great Atlantic Sargassum Belt reached a record abundance in 2025 and remained near that level in 2026, when the Caribbean Sea and Gulf of Mexico experienced record regional accumulations. Read only as measurements of biomass, these developments might appear to point in opposite directions. One concerns anxiety over the removal of marine organisms; the other concerns an apparently vigorous expansion of marine life.

Ecologically, however, they describe the same structural problem.

Antarctic krill remain abundant when measured across the immense area of the Southern Ocean. The 620,000-tonne fishery trigger represents a small fraction of estimated regional biomass. That aggregate comparison has consequently become central to arguments that the fishery remains sustainable. Yet whales, penguins, seals, fish, and seabirds do not consume a regional biomass estimate. They consume krill that must be present within reachable foraging distance, at the correct depth, during the correct season, and in sufficient density to repay the energy expended finding and capturing them.

A penguin feeding chicks cannot substitute krill located hundreds of kilometres away for krill removed near its colony. A recovering whale population cannot consume the portion of the statistical stock that is inaccessible during migration. Biomass may exist at the scale of the management model while failing to exist operationally at the scale of the animal.

The recent fishery closures sharpen this distinction because they followed a change in spatial governance. Until the end of the 2024 season, a conservation measure distributed the Area 48 catch among several subareas. When that measure lapsed, the aggregate trigger level remained, but the restrictions governing where the catch could be taken disappeared. The numerical ceiling did not rise. The geometry beneath it changed.

Fishing vessels could concentrate more effectively in krill-rich locations. Those locations are rich for the same reason that predators gather there: currents, bathymetry, seasonal production, and krill behaviour produce dense, accessible swarms. Removing spatial restrictions therefore increased the fleet's effective capability without increasing its nominal authorization. A limit designed at the scale of the fishery could remain intact while protection at the scale of the food web deteriorated.

The early closures do not, by themselves, prove that Antarctic krill have been depleted below a biologically sustainable global or regional level. They establish something more precise: industrial harvesting capacity can now traverse the permitted catch space faster and concentrate within it more

strongly. Whether that capability produces ecosystem damage depends on where and when extraction occurs, which predators share those locations, how rapidly krill are replenished, and whether climate-driven changes in sea ice and circulation are already narrowing the predators' alternatives.

This is an ecological admissibility problem. Krill can exist, be scientifically observed, enter a biomass estimate, and remain legally harvestable without remaining available to every organism whose continuation depends on them. Existence does not imply reachability. Reachability does not imply availability at the necessary time. Availability does not imply that extraction is recoverable. An aggregate stock can remain above a precautionary threshold while local food-web continuations become progressively harder to sustain.

Sargassum reveals the inverse failure.

Floating Sargassum is not inherently pollution. In the open ocean it forms a mobile habitat supporting fish, crabs, molluscs, microorganisms, seabirds, and sea turtles. It captures carbon through photosynthesis and creates shelter and feeding surfaces in waters that might otherwise possess relatively little structural complexity. Within the Sargasso Sea, its presence is part of an established ecological regime.

The Great Atlantic Sargassum Belt is a different spatial phenomenon. Since 2011, enormous seasonal aggregations have repeatedly formed across the tropical Atlantic, extending from West Africa toward the Caribbean and Gulf of Mexico. The causes remain under investigation and likely include changing circulation, winds, rainfall, nutrient delivery, river discharge, atmospheric deposition, upwelling, and warmer waters. No single cause adequately explains every year or region.

When Sargassum reaches the coast in extreme quantities, its ecological function can reverse. Dense mats reduce light reaching corals and seagrasses, obstruct the movement of animals, alter water circulation, and physically smother shallow habitats. Once stranded or trapped near shore, the accumulated material begins to decompose. Microbial decomposition consumes dissolved oxygen, while the decaying biomass releases hydrogen sulfide, ammonia, and other compounds. Material that served as habitat offshore becomes an anoxic and sometimes toxic burden at the coastline.

Nothing has disappeared during this transition. The biomass remains present, indeed, its excessive presence is the problem. What changes is the receiving system's capacity to incorporate it.

A beach, mangrove margin, seagrass bed, municipal collection service, or coastal decomposition zone can process some arrival of organic matter. Beyond a certain rate, however, incoming biomass exceeds the available oxygen, space, labour, transport, storage, and decomposition pathways. The surplus does not become additional ecological wealth merely because it contains carbon and nutrients. It becomes maladmitted biomass: biological material arriving in a quantity, location, or state for which the receiving system has no viable continuation.

The krill and Sargassum cases therefore cannot be understood by placing their tonnage into a common planetary account. One system may retain a large total stock while losing strategically located biomass. Another may gain tens of millions of tonnes while becoming less habitable where that biomass accumulates. A global calculation could record the continued existence of Antarctic krill and the proliferation of Atlantic Sargassum as evidence that marine biomass remains plentiful. Such a calculation would miss the actual failures occurring in both systems.

The distinction is not between “too little life” and “too much life.” It is between biomass as quantity and biomass as ecological capability.

Biomass becomes capability only through relations. It must be reachable by the organisms that use it, synchronized with their life cycles, compatible with the chemistry of the receiving environment, and supported by pathways that redistribute its matter after death. Its value depends on what can eat it, inhabit it, transform it, transport it, decompose it, or preserve it. The same tonne of biological material can function as food, shelter, carbon reservoir, industrial feedstock, oxygen-consuming waste, or toxic obstruction depending on its position within those relations.

This is why the search for an “optimal biomass level for Earth” cannot be reduced to maximizing a global total. Even if total planetary biomass could be estimated reliably, the number would conceal its composition and topology. A biosphere dominated by a small number of rapidly proliferating organisms might contain substantial biomass while possessing reduced biodiversity and resilience. A region of exceptional primary production might export its carbon efficiently into deep water, support a complex food web, accumulate as undecomposed waste, or collapse into hypoxia. Equal quantities do not imply equal ecological function.

Temperature enters this problem as one condition among many. A warmer Earth may increase growing seasons and biological production in some regions. It may also strengthen ocean stratification, reducing the movement of nutrients from deeper water into the sunlit surface. Some of the most productive marine environments are cold not because coldness alone creates biomass, but because cold-water circulation and upwelling maintain access to nutrients. Conversely, warm nutrient-rich waters can support extraordinary blooms whose eventual decomposition damages the systems around them.

The relevant question is therefore not whether warmth or cold produces more life in the abstract. It is which combinations of temperature, nutrients, circulation, oxygen, acidity, seasonality, and food-web structure preserve productive biomass as an available and renewable ecological relation.

This distinction also changes how geoengineering should be evaluated. An intervention that increases net primary production cannot be presumed beneficial without asking what grows, where it grows, which organisms can use it, and what happens when it dies. An intervention that reduces biomass in one region cannot be presumed harmful without asking whether it suppresses a pathological bloom or prevents an anoxic transition. Gross production, carbon fixation, standing stock, biodiversity, harvestable yield, and ecological continuity are different quantities. They may sometimes rise together, but no general rule requires them to do so.

The Caldera Reactor and its associated biomass systems begin from this problem of incorporation. Managed kelp farms would provide a relatively predictable marine feedstock grown within an infrastructure designed to receive and process it. Excess Sargassum would enter as a secondary, disturbance-responsive stream. The purpose would not be to reward or perpetuate anomalous blooms. It would be to intercept some fraction of the biomass before coastal deposition converts it from off-shore habitat into an oxygen-consuming burden.

The distinction between managed kelp and opportunistic Sargassum is essential. If the reactor became economically dependent on pathological Sargassum abundance, it could create an institutional interest in the continuation of the ecological failure it was meant to relieve. Surrounding kelp farms

instead provide a controllable baseline substrate, while the adaptive processing system offers additional capacity when anomalous biomass arrives. The infrastructure remains viable without requiring the disturbance.

Processing alone does not guarantee ecological benefit. Collection vessels expend energy. Wet biomass is expensive to transport and dewater. Sargassum may contain salt, sand, plastics, arsenic, heavy metals, or decomposed material that complicates conversion. Biocrude returns carbon to the atmosphere when burned, whereas durable bioplastics or stable carbon products may retain it longer. The relevant carbon balance depends on the complete pathway and on what would have happened to the Sargassum without interception.

Nevertheless, the conceptual role of the reactor is clear. It supplies an alternative continuation for biomass that a coastal system cannot incorporate. It changes the material's admissible future: from uncontrolled deposition and decomposition toward differentiated products, recoverable energy, durable carbon forms, or carefully managed residues. Its success would be measured not by the gross amount of biomass consumed but by the ecological damage avoided, useful displacement achieved, and carbon retained.

Krill management poses the complementary task. There the objective is not to create an alternative use for surplus biomass, but to preserve local availability within an existing food web. Spatial catch constraints may appear economically inefficient because they prevent vessels from concentrating where harvesting is easiest. Ecologically, that inefficiency is precisely their function. By dispersing extraction, they preserve distinctions among feeding regions and reduce the probability that industrial capability will collapse a locally essential swarm into an aggregate statistic.

The two cases reveal a general principle:

Biospheric health depends not on how much living material exists, but on whether biomass remains appropriately distributed, reachable, transformable, and incorporable across the systems whose continuation depends upon it.

This principle will govern the interventions considered throughout this book. Polar ice will not be evaluated merely by tonnes frozen, but by reflective area, seasonal persistence, ecological function, and the larger ice regimes thereby preserved. Rainforest generators will not be evaluated merely by vertical biomass, but by habitat complexity, nutrient continuity, maintenance burden, and increasing ecological autonomy. Kelp cultivation will not be evaluated solely by growth rate, but by its effects on marine food webs, carbon pathways, oxygen, harvest cycles, and material use. Geoengineering will not be treated as the production of desirable quantities in isolation. It will be treated as the design of relations through which those quantities acquire, or lose, the capacity to sustain a world.

Chapter 3

Climate Change as Commitment Mismatch

Sea-level rise draws a disproportionate share of public and political attention among the consequences of climate change, and for good reason. It is visible, cumulative, and largely irreversible on human timescales, and it threatens some of the most densely populated and economically valuable land on Earth. But the reason sea-level rise feels distinctively urgent is not that it is a separate category of harm from everything else climate change produces. It is that it makes unusually legible a structural problem that runs through nearly every consequence of a changing climate: human and ecological systems become committed to a set of climatic regularities, often over decades or centuries, and climate change revises those regularities after the commitment has become expensive or impossible to undo.

Commitment, in this sense, is not a choice made carelessly. It is the ordinary and often rational outcome of building anything durable. A coastal city is not built in a single year by people ignorant of sea level; it accumulates across generations, each addition assuming that the shoreline, tidal range, and storm-surge frequency observed at the time of construction will remain approximately stable. A seawall calibrated to a particular flood return period, a subway system built below a particular water table, a port whose depth and layout assume a particular sea level, and the property values, tax base, and insurance markets that depend on all three, together constitute a vast, distributed bet that the coastline of the recent past will resemble the coastline of the future. That bet was reasonable for most of the period during which it was made. It has become progressively less reasonable without any single decision-maker having chosen to make it so.

The same structure appears wherever commitment and climatic regularity diverge, whether or not the result resembles a coastline.

Agricultural regions commit to a rainfall pattern through crop selection, irrigation infrastructure, land tenure, and the accumulated knowledge of farmers about when to plant and what to expect. A region that has reliably received a wet season followed by a dry season will develop crop varieties, storage systems, and credit cycles suited to that pattern, and those investments are not easily redirected when the pattern shifts. A multi-year drought does not simply reduce yield in the years it occurs; it can exhaust the assumptions embedded in seed stock, debt schedules, and water rights that were calibrated to a climate the region no longer has.

Water utilities in mountainous and high-latitude regions commit to a snowpack that accumulates in winter and releases meltwater gradually through spring and summer, smoothing supply across a season that would otherwise alternate between flood and drought. Reservoir capacity, treatment infrastructure, and downstream agreements among competing users are sized to this seasonal release.

When warming shifts precipitation from snow to rain, or accelerates snowmelt into a shorter and earlier pulse, the physical volume of water available across a year may change relatively little while its timing changes enough to overwhelm infrastructure built for a different release curve.

Coral reefs commit, in an ecological rather than an institutional sense, to a narrow band of sea-surface temperature within which the symbiotic algae that reefs depend on for energy and color can survive. This commitment is written into the reef's own biology rather than into concrete or contracts, but it produces the same structural vulnerability: a system optimized for a historical regularity has no graceful response available when that regularity shifts faster than the system's own capacity for adaptation or migration.

Insurance and reinsurance markets commit to actuarial tables built from historical loss data, on the assumption that the frequency and severity of future events will resemble the frequency and severity of past ones closely enough for the tables to remain solvent. Fisheries commit to seasonal ice cover, migration timing, or spawning-ground temperature. Disease vectors commit, and human populations commit around them, to a range within which a given mosquito, tick, or pathogen can or cannot survive a given latitude's winter. In each case, a system, whether institutional, ecological, or actuarial, has organized itself around a regularity that persisted long enough to be treated as a stable feature of the world, and climate change is now revising that regularity on a timescale shorter than the system's ability to reorganize around a new one.

This framing has several consequences for how climate risk should be assessed and how interventions should be evaluated.

First, it explains why the same magnitude of physical change can produce wildly different amounts of harm depending on what has been built around the prior condition. A degree of warming that a sparsely settled, agriculturally flexible region absorbs with modest adjustment can devastate a densely built, heavily indebted, institutionally rigid region experiencing the identical physical change. Vulnerability is not a direct function of exposure to a hazard; it is a function of exposure combined with the depth and inflexibility of prior commitment. This is why coastal megacities, snowpack-dependent water systems, and narrow-tolerance ecosystems recur as the most frequently cited examples of climate risk: not because their physical exposure is uniquely severe in an abstract sense, but because their commitments are unusually deep, unusually expensive to relocate, and unusually slow to renegotiate.

Second, it clarifies why adaptation is not simply a matter of technical capacity but of timing. A commitment can in principle be revised, a seawall raised, a crop variety changed, a reservoir rebuilt, a reinsurance table repriced, but revision takes time, capital, and institutional coordination that themselves have a rate. If the underlying climatic regularity changes faster than the commitment can be revised, the system experiences the change as a shock even though the physical trend that produced it may have been gradual and, in principle, foreseeable for decades. Much of what appears as sudden climate catastrophe is better described as the delayed collision between a slow physical trend and a commitment whose revision rate was slower still.

Third, and most directly relevant to the chapters that follow, this framing reorients what a geoengineering intervention should be evaluated against. An intervention should not be judged primarily by whether it returns some global variable toward a historical value. It should be judged by whether it preserves, extends, or safely revises the reachability of the specific commitments that populations

and ecosystems have already made, and by whether it does so without creating new commitments whose own revision, should the intervention need to be altered or withdrawn, would prove equally rigid. An intervention that stabilizes a coastline while making the underlying climatic trend worse elsewhere has not resolved a commitment mismatch; it has relocated one. An intervention that succeeds only so long as it continues indefinitely has not restored a lost regularity; it has substituted a new commitment, to the intervention itself, for the old one, and that substitution carries its own risks if the intervention is later withdrawn, defunded, or rendered obsolete.

This last point anticipates the central concept of the following chapter. An intervention that a system comes to depend on cannot simply be removed once it has been in place long enough for other systems to reorganize around it. Stopping it does not restore the world that existed before it began; it produces a new transformation, potentially more disruptive than either the original problem or the intervention's continued operation. This is termination shock in its general form, and it is one of the central risks that continuation geometry, developed next, is built to make explicit and evaluable before any intervention proposed in this book is undertaken at scale.

Chapter 4

Continuation Geometry

The preceding chapters have used several words informally that now need to be given exact meaning: reachable, viable, recoverable, committed. Chapter 3 leaned on all four without defining any of them, because the argument at that stage needed only their intuitive sense. Formalizing them here is not a detour into abstraction for its own sake. It is what allows every intervention discussed in the rest of this book, from a nuclear-powered ice system to a routed iceberg to a planetary infrastructure network, to be evaluated against a single, consistent standard rather than assessed impressionistically case by case.

The central idea is that an intervention does not simply move a system from one state to another. It changes which later states remain available to reach. A climate intervention, a piece of infrastructure, or an ecological disturbance can be described not only by where it puts a system today, but by the shape of the space of futures that system can subsequently move through. This chapter calls that space, and its evolution under intervention, continuation geometry.

Reachability

A state is reachable from a given starting point if some sequence of physically and institutionally permitted transitions connects the two. Reachability is the weakest of the concepts developed in this chapter, and it is important to establish it first precisely because it is weak: many proposals are defended by showing that a desired outcome is reachable in principle, without showing that it is reachable under the constraints that actually govern the system, or that reaching it would leave the system able to do anything else afterward.

Consider polar sea ice. It is reachable, in the narrow sense, to freeze a given volume of seawater given sufficient energy and time; nothing in the physics forbids it. But this observation alone says nothing about whether the resulting ice performs any useful climatic function, whether the energy expenditure and waste heat involved are compatible with the surrounding system's other commitments, or whether the frozen state persists long enough to matter. Reachability answers the question "can this state be reached at all." It does not answer "should it be," "can it be reached without foreclosing other states," or "can it be maintained once reached." Those questions require the stronger concepts that follow.

Viability

A state, or more precisely a trajectory through a sequence of states, is viable if the system can occupy it while continuing to satisfy the constraints that define its ongoing function. A power grid is not viable merely because it can be made to deliver electricity for a single afternoon; it is viable if it can deliver electricity continuously while maintaining frequency, voltage, and reserve margins within the tolerances its connected systems require. A coral reef is not viable merely because individual corals can survive a single marine heatwave; it is viable if the reef's net calcification, recruitment, and recovery processes can keep pace with disturbance over the timescales on which the reef must persist.

Viability is therefore a claim about a region of state space, not a single point within it. Engineers and ecologists both use versions of this concept, often called an admissible region or a safe operating envelope, to describe the set of conditions under which a system continues to function rather than degrading, collapsing, or requiring outside rescue. An intervention can move a system into a reachable state that lies outside its viable region: the state exists, the system can be pushed into it, but the system cannot remain there and continue doing what it needs to do. This is the failure mode that a purely reachability-based evaluation misses and that a viability-based evaluation is designed to catch.

Viability regions are not fixed. They shift as a system's surrounding conditions change, and this is precisely what climate change does to the systems described in Chapter 3. A coastal city's viable region, defined by the combination of sea level, storm frequency, and infrastructure capacity within which it can continue to function without catastrophic loss, contracts as sea level rises, even if no single event pushes the city outside it. The commitment mismatch discussed in the previous chapter can now be restated more precisely: a system's commitments were made on the assumption of a particular viable region, and climate change is shrinking or relocating that region faster than the system can revise its commitments to match.

Recoverability

A state is recoverable if, once a system enters it, some available and sufficiently fast process can return the system to its viable region before irreversible loss occurs. Recoverability is what separates a viability excursion, a temporary departure from the admissible region that resolves before permanent damage accrues, from a genuine collapse.

Recoverability depends on rate as much as on mechanism. A drought that pushes a rain-fed agricultural region outside its viable moisture range is recoverable if the region has water reserves, alternative crops, or credit sufficient to bridge the gap until rainfall resumes. The same drought, extended long enough, exhausts those reserves and crosses into unrecoverable territory: soil degrades past the point of restoration within a farming generation, seed stock is consumed rather than replanted, debt forecloses land that would otherwise have supported recovery. The physical hazard did not change in kind between the recoverable and unrecoverable cases; what changed was whether a sufficiently fast countervailing process remained available.

This gives recoverability a strict boundary rather than a fuzzy one, at least in principle: for a given system, there exists some maximum duration and severity of excursion beyond which no available

process returns it to viability in time. Locating that boundary precisely is often difficult, since it depends on reserves, institutional capacity, and processes that are only partially observable in advance. But the existence of the boundary, and the fact that it can in principle be estimated and monitored, is what allows recoverability to function as an engineering standard rather than merely a retrospective judgment applied after a collapse has already occurred.

Path Dependence

The three concepts above describe a system's relationship to a given state or trajectory at a given time. Path dependence describes how that relationship changes as a consequence of the route by which the system arrived. Two systems occupying what appears to be the identical present state can have different viable regions and different recoverability boundaries going forward, because the route each took to reach that state altered capacities, dependencies, or reserves that are not visible in the state description alone.

A reservoir at fifty percent capacity that arrived there through a single dry season behaves differently, going forward, than a reservoir at fifty percent capacity that arrived there through a decade of gradually declining snowpack, because the second reservoir's watershed, downstream ecosystems, and water-rights agreements have already begun reorganizing around a lower baseline in ways the first reservoir's have not. The present state is the same number; the continuation geometry is not.

Path dependence is what makes intervention timing, not only intervention magnitude, a first-order variable throughout this book. An intervention introduced early, while a system's viable region is still close to its historical range, can often be modest in scale and still effective. The identical intervention introduced later, after the system's surrounding dependencies have reorganized around a degraded baseline, may need to be far larger to achieve the same effect, or may no longer be sufficient at any scale, because the system's recoverability boundary has itself moved.

Termination Shock as the Paradigmatic Case

These four concepts converge most sharply in the phenomenon that gives Neal Stephenson's novel its title, and that this book has referenced only informally until now. An intervention is introduced into a system. The system's surrounding dependencies, infrastructure, agriculture, institutions, and in some cases entire ecosystems, begin adapting to the intervention's presence, not to the prior natural state it was designed to approximate. This adaptation is itself a form of path dependence: the system's continuation geometry reorganizes around the intervention as a standing feature of its environment rather than as a temporary correction.

Once this reorganization has occurred, removing the intervention does not restore the system to the state that existed before the intervention began. It produces a new transition, from a state adapted to the intervention's presence to a state that must suddenly readapt to its absence, and this transition can be more severe, and can cross recoverability boundaries the original problem never approached, than either the initial climatic disturbance or the intervention's continued operation would have.

This is why termination shock cannot be treated as an unfortunate side effect to be managed after the fact. It is a predictable consequence of path dependence applied to any intervention sustained long

enough for dependent systems to reorganize around it, and it means that every intervention discussed in this book must be evaluated not only by its effect while operating, but by the reachability, viability, and recoverability of the transition that would follow its slowing, failure, or deliberate withdrawal. An intervention whose benefits depend on indefinite continuation has not solved a commitment mismatch. It has substituted a new commitment, to itself, for the one it was meant to relieve, and that substitution must be priced into the intervention's design from the outset rather than discovered during an eventual termination.

Continuation Geometry as an Evaluative Standard

Taken together, reachability, viability, recoverability, and path dependence supply the standard the rest of this book applies. A proposed intervention will be asked: what states does it make reachable that were not reachable before, and at what cost to states that were reachable and now are not. What is its effect on the viable region of the systems it touches, including systems it was not explicitly designed to affect. What is the recoverability boundary of the trajectories it creates, and how does that boundary compare to the boundary of the trajectory it replaces. And what does the intervention's own termination look like, given the path-dependent reorganization its continued operation would produce.

This standard does not by itself tell us which interventions to pursue. It tells us what a defensible answer to that question must account for, and it rules out the two simplest but inadequate answers this book has already rejected: that an intervention is justified because it moves a single planetary variable in a preferred direction, and that an intervention is justified because it increases some aggregate quantity, such as biomass, without regard to the geometry of the commitments built around it. Part II now turns to the practical problem this standard creates: how competing interventions, each with its own effect on reachability, viability, recoverability, and path dependence, can be explored, compared, and contested before any of them is committed to at planetary scale.

Part II

Seeing Before Acting

Chapter 6

Simulation Before Commitment

Chapter 4 closed with a standard: an intervention should be evaluated by what it makes reachable, what it removes from the viable region of the systems it touches, how recoverable its trajectory is, and what its own termination would look like given the path dependence its operation creates. That standard presupposes something this book has not yet supplied: a means of applying it *before* an intervention is built, rather than after its consequences have already begun to accumulate. A continuation-geometric evaluation performed retrospectively, once aerosols are already in the stratosphere or a reactor is already drawing seawater, is a post-mortem. This chapter and the one that follows ask what it would take to perform that evaluation prospectively, across competing proposals, before any of them is committed to at planetary scale.

The obstacle is not technical in the first instance. It is institutional, and it has a specific shape: the mechanisms modern societies use to decide what large collective actions to take remain organized around narrative coordination, discourse, persuasion, and symbolic alignment carried out against a background that is presumed, largely, to stay fixed while the argument proceeds. A debate about whether to pursue polar albedo maintenance, oceanic iceberg routing, or a planetary material-circulation network resembles, in its institutional form, a debate about a tax policy or a zoning variance: position papers, hearings, coalitions, votes. But the systems these interventions touch, ice sheets, ocean circulation, food webs, are exactly the high-dimensional, nonlinear, strongly path-dependent systems for which narrative coordination was never well suited, because narrative form survives only by suppressing the dimensionality that determines what will actually happen.

Where Narrative Coordination Fails

Narrative coordination works by compressing a possibility space down to a small number of salient variables arranged in causal sequence: a policy is proposed, a mechanism is asserted, a consequence is claimed, and the argument is won or lost on the plausibility of that chain as told. This is not a defect introduced by bad faith or insufficient rigor. It is intrinsic to the medium. A story that preserved every relevant dimension of a coupled ocean-atmosphere-ecosystem-economic system would not be a story anyone could tell or follow; it would already be a simulation. Narrative coordination therefore functions well exactly where Chapters 1 through 3 showed the geoengineering imagination tends to go wrong: toward a single target variable, a single mechanism, a single preferred future, because those are the forms a story can hold.

This is why the debates that precede planetary interventions so often resemble the debates Chapter

1 diagnosed in the temperature case and Chapter 2 diagnosed in the biomass case. A proposal is defended by a story about what it will do to one variable, global temperature, total carbon sequestered, tonnes of ice restored, and the story's persuasiveness is judged by its internal coherence rather than by whether the variable it names is the one that mattered. The commitment-mismatch systems of Chapter 3, coastal infrastructure, rain-fed agriculture, snowpack-dependent water utilities, do not appear in such a story except as later casualties, because a narrative sufficient to hold all of them simultaneously, in their actual coupling to one another and to the proposed intervention, is not a narrative any single speaker can construct or any single audience can evaluate by listening.

Nor does narrative coordination expose termination. A proposal is typically argued, won, and funded on the strength of its operating case; its withdrawal case, what happens to the systems that will have reorganized around it if it must later be slowed, defunded, or reversed, is rarely part of the story that secures its adoption, because a story about ending well is a weaker story than a story about beginning well, and narrative coordination rewards the stronger story rather than the more complete one. Stephenson's Schmidt succeeds narratively before he succeeds physically: the account of what his installation will do to a single variable, the sea-air temperature difference driving hurricane intensity, is compelling enough to mobilize capital and attention long before anyone outside his organization has modeled what its slowing or failure would do to the monsoon systems and agricultural calendars that will have reorganized around its operation. This is not a failure of Schmidt's intelligence. It is a demonstration of what narrative coordination structurally cannot represent.

What Simulation Already Does Elsewhere

The relevant contrast is not hypothetical. Contemporary industrial production has already largely abandoned narrative coordination in favor of a different epistemic mode. Aerospace assembly, semiconductor fabrication, and large energy systems are not designed by debating a story about what a proposed configuration will do; they are designed by instantiating the configuration inside a digital twin or factory simulator and observing what a full-dimensional model of throughput, latency, failure mode, and resource interdependence actually does under the constraints that will govern it in operation. Feasibility is established before material commitment, not argued toward it. The narrative account that eventually accompanies the decision, delivered to investors, regulators, or the public, is typically reconstructed after the simulation has already determined the outcome; it legitimizes a decision the model has made rather than making the decision itself.

This is not a claim that industrial simulation is infallible, only that it represents a different and, for high-dimensional path-dependent problems, structurally more adequate way of relating belief to commitment. A story can be revised costlessly when it loses persuasive force; a simulated trajectory, once explored, has already exposed whatever it exposed, and that exposure cannot be argued away. The asymmetry matters because it is precisely the asymmetry continuation geometry needs. A viability region, a recoverability boundary, a path-dependent reorganization around a sustained intervention, these are properties of a trajectory through a high-dimensional state space, and they are visible to a model that propagates that trajectory forward under constraint. They are not reliably visible to a narrative, because a narrative is, definitionally, a compression that has already discarded most of the dimensions along which such properties would show themselves.

The geoengineering proposals this book will evaluate in Parts III through VI, active polar albedo

maintenance, oceanic continuation engineering, constructed terrestrial productivity, planetary material circulation, are precisely the kind of coupled, irreversible, long-horizon interventions for which the industrial precedent is most directly applicable and for which narrative evaluation is least adequate. A proposal to route icebergs as mobile boundary conditions for regional albedo, or to reintroduce Sargassum-derived carbon into a terrestrial productivity system through a caldera-reactor pathway, cannot be responsibly evaluated by the story of what it will do to one number. It can be evaluated by observing what a model of the coupled system it enters actually does when the intervention is instantiated inside it, perturbed, and run forward far enough to expose whether the resulting trajectory is viable, recoverable, and safely terminable.

Preference Revealed Rather Than Declared

A further consequence follows from adopting simulation rather than narrative as the substrate for evaluating these interventions, and it bears directly on the question Chapter 1 raised and refused to answer: what future should be preferred, given that no single scalar captures planetary health. Narrative coordination treats preference as something that precedes deliberation and can be declared: a constituency wants a stable coastline, a lower global mean temperature, a preserved fishery, and the political process aggregates these declared wants through voting, negotiation, or rhetorical victory. But Chapters 1 and 2 have already shown that declared preferences of this kind, expressed as scalars, obscure exactly the tradeoffs a genuine evaluation needs to surface. A constituency that declares a preference for a lower global temperature has not thereby specified whether it prefers the upwelling-dependent marine productivity of a colder ocean or the extended growing seasons of a warmer one, because that specification requires confronting a tradeoff the declared preference was never asked to resolve.

A simulation-based civic substrate handles this differently. Rather than asking a population to declare, in advance, what tradeoff it prefers among biomass, biodiversity, coastal stability, and agricultural productivity, it allows the population to explore trajectories in which those quantities move together and against one another under an explicit set of physical and ecological constraints, and to discover, through sustained interaction with those trajectories, which tradeoffs it is actually willing to accept. Preference, on this view, is not an antecedent fact waiting to be reported; it is a functional over trajectories, revealed by which paths a population continues to inhabit, develop, and defend once their consequences are no longer hypothetical but explored. This is not a claim that simulation manufactures consensus. Populations exploring the same coupled model can and will diverge, some sustaining a trajectory that accepts reduced upwelling-driven marine productivity in exchange for extended agricultural range at temperate latitudes, others refusing that trade. The simulation does not resolve this disagreement. It relocates the disagreement from the level of slogan, warmer is better, colder is safer, to the level of structure: which specific tradeoff, under which specific constraint, is being accepted or refused, and by whom.

This relocation is what makes the standard from Chapter 4 usable as a civic instrument rather than only an analyst's checklist. Reachability, viability, recoverability, and path dependence are properties that a population can be shown, inside a shared simulated environment, to be trading against one another as they navigate toward a preferred trajectory, rather than properties an expert asserts on their behalf after the fact. A community exploring an iceberg-routing proposal for a particular coastline can observe, within the simulation, not only whether the proposal stabilizes the coastline it targets

but what the routing does to current patterns and fishery yields elsewhere, and can revise its own preference once that coupling becomes visible rather than remaining an abstraction. Legitimacy, on this account, does not derive from having voted for a slogan. It derives from having navigated, with open eyes, the structure that a chosen trajectory actually imposes.

Irreversibility as the Design Constraint

Two features of the industrial precedent require modification before it can serve a civic purpose, and both concern irreversibility. First, a factory simulator's trajectories are, in an important sense, cheap: a failed configuration can be discarded before a single component is manufactured, and no cost attaches to having explored it. A civic simulation intended to evaluate planetary interventions cannot offer this cheapness without forfeiting exactly the discipline that makes simulation preferable to narrative in the first place. If every explored trajectory can be costlessly abandoned, exploration degenerates back into narrative, a proliferation of stories about what might happen, none of them binding. The architecture developed in the next chapter, adapted from an existing framework for treating civilizations as executable hypotheses, addresses this by distinguishing exploratory branches, which may be freely abandoned, from committed trajectories, which accumulate history and cannot be costlessly rolled back once a population has begun to act on them. Exploration remains free; commitment does not.

Second, and more specific to the geoengineering case, the trajectories a civic simulation must expose are not only trajectories of operation but trajectories of termination. An industrial digital twin is rarely asked to model what happens if the factory it represents is shut down after thirty years of continuous operation, because a factory's surrounding dependencies do not typically reorganize around its indefinite continuation in the way a regional climate, an agricultural calendar, or a coastal ecosystem can reorganize around a sustained geoengineering intervention. Chapter 4 identified this reorganization as the mechanism of termination shock. A civic simulation adequate to geoengineering must therefore model not only the trajectory an intervention creates while operating, but the trajectory its slowing, failure, or deliberate withdrawal would create given whatever reorganization its operation has already produced, and must expose that withdrawal trajectory to the same population evaluating the intervention's adoption, before adoption rather than after.

From Infrastructure to Interventions

None of this yet specifies what such a simulation would look like in enough detail to be built, governed, or trusted. That specification is the task of the next chapter, which takes up a second existing framework, one that treats civilizations themselves as executable hypotheses rather than settled arrangements, and asks how its architecture of layered admissibility, physical, institutional, and narrative, and its machinery of repair, model revision, institutional correction, infrastructure restoration, can be adapted specifically to the problem this book has posed: comparing polar, oceanic, terrestrial, and planetary-circulation interventions against one another and against their own termination, before any of them is built. What this chapter has established is only the negative case and the positive precedent: that narrative coordination is structurally unsuited to evaluating interventions of this kind, and that a mode of coordination already exists, in industrial production, that is not. The remaining question is how to civicize it.

Chapter 7

Civilizations as Executable Hypotheses

Chapter 6 established the negative case: narrative coordination is structurally unsuited to evaluating interventions whose consequences are high-dimensional, nonlinear, and path-dependent, and an alternative epistemic mode, already dominant in industrial production, exists that is not. It also left two demands unmet. A civic simulation adequate to geoengineering must distinguish exploration, which should remain free, from commitment, which should not; and it must expose termination trajectories alongside operating trajectories, before either is chosen. This chapter draws on an existing framework, developed independently under the name *Civilizations as Hypotheses*, that supplies an architecture meeting both demands, and adapts it specifically to the problem of comparing polar, oceanic, terrestrial, and planetary-circulation interventions against one another and against their own withdrawal.

A Civilization as a Parameterized Hypothesis

The framework's foundational move is to stop treating a civilization, or any of its major infrastructural commitments, as a settled arrangement to be described, and start treating it as a hypothesis to be tested. Formally, a civilization is represented as a tuple

$$C = (A, R, P, N, \Sigma),$$

where A denotes the assumptions a civilization makes about behavior, incentives, and institutional response; R the formal rules governing permissible state transitions; P the preference fields, treated as dynamic rather than fixed, describing how collective wants evolve under experience and counterfactual exposure; N the norms, the commitments a civilization treats as inviolable regardless of the optimization pressure bearing on them; and Σ the informational feedback the simulator itself exerts on the reflective states of the agents inhabiting it, since a population that can observe the consequences of its own trajectory is not the same population, epistemically, as one that cannot.

This formalization matters for the present book because it gives content to a distinction Chapter 4 needed but could not yet supply: the distinction between a proposed intervention's physical possibility and its actual admissibility once institutional and normative constraints are taken into account. A civilization is not a doctrine to be affirmed or rejected as a whole. It is a parameterized object whose components can be independently varied, tested under perturbation, and revised in light of failure, exactly as any of the interventions in Parts III through VI must be.

The Civic Simulator: What It Is and Is Not

The instrument through which such a hypothesis is tested is called, in the source framework, the civic simulator: a persistent, multi-model computational environment in which competing formalizations of A , R , P , and N are instantiated as executable models, subjected to perturbation, and compared. Three existing institutions supply the right analogies, and each clarifies what the simulator does by clarifying what it does not do. A wind tunnel does not tell an engineer what shape to build; it tells the engineer what happens to a candidate shape under load. A flight simulator does not determine where a pilot should fly; it determines whether the pilot has the skill to survive turbulence en route. A trial registry does not decide which drug to approve; it determines which predictions were registered in advance and which of them held. The civic simulator, applied to planetary intervention, does not predict which future Earth will have, and it does not decide which intervention should be built. It reveals the structure of disagreement about which intervention should be built, by making the disagreement's actual load-bearing assumptions visible and separately testable.

This is a stronger claim than it may first appear, because it forecloses two temptations this book has already warned against. It forecloses treating the simulator as an oracle that returns a single correct planetary policy, since Chapter 1 has already shown that no such single policy exists to be returned. And it forecloses treating the simulator as a mere visualization aid for a decision that will actually be made elsewhere, by narrative or by fiat, since a visualization aid carries no consequence and therefore reproduces exactly the cost-free revisability that Chapter 6 identified as narrative coordination's central defect.

Stratified Admissibility

A proposed intervention, an iceberg routed to stabilize a regional albedo, a caldera reactor processing surplus Sargassum, a network of orthodromic distribution reserves, does not face a single admissibility test. It faces three, nested rather than merged. Physical admissibility, A_P , asks whether the intervention satisfies the thermodynamic, material, and engineering constraints that govern it regardless of any human decision: whether the ice can be routed at all given ocean currents and structural integrity, whether the reactor's energy balance closes, whether the distribution network's capacity meets its demand. Institutional admissibility, A_I , asks whether the intervention satisfies the regulatory, jurisdictional, and coordination constraints layered on top of physical possibility: whether the routing crosses waters under contested sovereignty, whether the reactor's emissions fall inside an existing permitting regime, whether the network's siting decisions survive the placement objective developed in Chapter 5. Narrative admissibility, A_N , asks whether the intervention is permitted by prevailing public belief and the cultural resistance or acceptance a population brings to interventions of its kind, independent of whether it is physically or institutionally sound.

These three regions are nested, and the gap between them is diagnostic rather than incidental. An intervention can be physically executable, A_P , while remaining institutionally stalled, outside A_I , and a great deal of what looks like technological failure in the geoengineering literature is in fact this gap misclassified as physical impossibility. This misclassification matters directly to the interventions this book will evaluate. A proposal for active polar albedo maintenance may fail not because the underlying thermodynamics are unsound but because no existing institutional framework can authorize action across the jurisdictions the intervention would need to cross; a proposal for planetary

material circulation may fail not because the placement mathematics of Chapter 5 cannot be solved but because no governance structure exists to bind the communities whose cooperation the solution requires. The civic simulator's role, in each case, is to keep these three admissibility regions visibly distinct, so that a genuine institutional or narrative obstacle is never quietly absorbed into, and thereby mistaken for, a physical one.

The Explanatory Stack

The framework organizes evaluation of any candidate intervention through a five-step derivation chain, and mapping this chain onto the present book's interventions clarifies what work remains for Parts III through VI. Step one, resource geometry, denoted (Φ, v, S) , asks what fundamentally constrains the system: available flux, velocity of transport, and storage capacity, the same quantities Chapter 5's placement model already made explicit for material circulation and that Parts III through V must make explicit in turn for ice, ocean, and terrestrial biomass. Step two, admissibility, denoted A , asks which futures survive the stratified test just described. Step three, repair, denoted R , asks how a future that fails admissibility can be returned to it, the machinery taken up in the next section. Step four, persistence, denoted P , asks how historical continuity is maintained across repair events, since a civilization, like the systems described in Chapter 3, cannot simply reset to zero after a failure without erasing the very commitments whose mismatch produced the failure. Step five, collective choice, denoted Φ_C , asks how decisions are actually formed, not through discrete votes on discrete slogans but through the aggregated gradient of revealed preference Chapter 6 already described.

This chain gives the book's remaining parts a shared internal structure. Each intervention evaluated in Parts III through VI can, and should, be run through the same five steps: what resource geometry constrains it, which admissibility region it actually occupies, what repair its failure modes require, what persistence its operation and eventual termination demand, and how the collective choice to pursue it would actually be formed rather than merely declared.

Reachability Volume as the Objective

The framework supplies a formal objective that generalizes, rather than replaces, the continuation-geometric standard of Chapter 4: maintain and expand the volume of navigable, admissible futures, $\text{Vol}_R(A) > 0$, rather than maximize any single scalar output. An intervention that stabilizes a coastline while foreclosing the admissibility region available to the fisheries and current patterns elsewhere is not, on this objective, unambiguously successful merely because the coastline metric improved; it must be asked whether the admissible volume of the coupled system, taken as a whole, expanded or contracted. A civilization, or an intervention, that maximizes present benefit while stripping the regenerative and repair capacity that sustains future admissibility is failing on this objective regardless of how favorably its present-tense metrics read, exactly as a coastal city that maximizes short-term development while exhausting its adaptive capacity is, in the terms of Chapter 3, deepening a commitment mismatch rather than resolving one.

This reformulates the question Chapter 1 posed and declined to answer with a scalar. The relevant question was never which single future to select. It is which futures remain reachable after a given intervention's choices have been made, and an intervention should be judged by whether it expands

or contracts that reachable volume, not by whether it moves one preferred coordinate toward one preferred value.

Disagreement Localization

Debates over geoengineering interventions typically present as binary: a proposal succeeds, or it fails; nuclear-powered ice maintenance is viable, or it collapses under its own cost and complexity. The framework's disagreement-localization procedure treats such binaries as underspecified rather than as genuine disputes, and resolves them by holding shared assumptions fixed and isolating the smaller number of testable parameters actually driving the divergence. Two models disagreeing about whether a given energy technology succeeds by a target date typically share far more than they dispute, agreeing, for instance, on reactor output or grid capacity, while diverging on a handful of specific, empirically testable parameters, such as regulatory timeline assumptions or the provenance of an assumed cost-learning curve. Once isolated, these parameters can be tested or monitored directly, rather than left to be litigated as an unresolvable clash of worldviews.

Applied to this book's material, this procedure is what separates a genuine dispute about, say, whether active polar albedo maintenance is admissible from a dispute that merely looks like one. Two positions on such a proposal may share every physical assumption in Part III and diverge only on an institutional parameter, whether cross-jurisdictional authorization can plausibly be secured within the intervention's required timeline, or a narrative one, whether affected populations will sustain support once early costs are visible. Localizing the disagreement to that parameter does not resolve the dispute, but it prevents the dispute from being fought, unproductively, on ground, thermodynamic feasibility, where no real disagreement exists.

The Machinery of Repair

A model that only detects admissibility failure without providing for its correction returns this book to the diagnostic-without-remedy position Chapter 3 was written to move past. The framework specifies four distinct repair channels, and distinguishing them matters because they operate on different timescales and require different institutional capacities. Infrastructure repair restores physical capacity directly, rebuilding grid capacity or ecological sinks, and is resource-intensive and time-bound in ways that make it the most legible but also the slowest channel. Institutional repair updates misaligned regulatory frameworks and rebuilds enforcement capacity, addressing failures that occur in A_I rather than A_P . Model repair identifies the specific assumptions generating predictive error and iterates them, maintaining a public, auditable provenance record of each revision, so that a model's history of correction remains inspectable rather than quietly discarded. Preference repair addresses distortion directly, providing populations whose option space has been artificially narrowed, by manipulation, deprivation, or the kind of narrative compression Chapter 6 described, with the experiences and counterfactuals needed to recognize that narrowing.

A repair efficiency measure, η_R , tracks the precise distance a disturbed trajectory is pulled back toward the admissibility boundary by a given repair action, giving repair the same quantitative standing as failure rather than treating it as an afterthought. This channel structure anticipates directly what Part VII will need: a geoengineering intervention that fails is rarely failing along only one channel,

and an adequate response distinguishes whether the failure is physical, institutional, model-level, or preference-level before proposing a remedy, since applying infrastructure repair to what is actually an institutional failure, or model repair to what is actually a preference distortion, wastes the resource-intensive channels on problems they cannot fix.

Testing Frameworks, Not Only Interventions

The framework's multi-agent configuration extends evaluation one level further, from testing individual interventions to testing the value systems used to select among them. Six idealized agent types, Optimization, which accepts any tradeoff that maximizes a stated utility; Stability, which penalizes deviation from a reference trajectory; Exploration, which maximizes the expansion of reachable option space; Conservation, which imposes hard constraints on physical throughput prior to any optimization; Innovation, which defers near-term constraints in expectation of future technological substitution; and Repair, which prioritizes restoration of admissibility above every other metric, are run simultaneously against the same coupled model and the same unexpected shocks. The question this comparison answers is not which ideology wins an argument, but which framework generates futures that remain navigable when perturbed in ways none of the agents anticipated.

This is directly usable by the parts that follow. The interventions of Parts III through VI, polar albedo maintenance, oceanic continuation engineering, constructed terrestrial productivity, planetary material circulation, are not merely technical proposals; each embodies, whether its proponents state this or not, a preference among Optimization, Conservation, Innovation, and Repair as governing logics. Running each intervention against all six agent postures, rather than presenting it under the framing its own proponents prefer, exposes which of its claimed virtues survive contact with a shock the intervention's own designers did not anticipate, which is precisely the exposure a narrative defense of the intervention would never produce.

Persistence Without Reset

The framework's final requirement returns directly to termination shock. A civic simulation that permits participants to reset a failed trajectory to zero, discarding its history and beginning again without cost, models a consequence-free reality that does not exist, and Chapter 3 has already shown at length that the actual world these interventions must operate in is one of accumulated commitment, not one of costless restarts. Persistence is required for three reasons the framework makes explicit: accountability, since agents must own the downstream consequences of interventions they propose rather than externalize them onto a reset function; institutional learning, since organizations accumulate tacit knowledge and memory only across a continuous history, not across repeated fresh starts; and preference evolution, since distinguishing a population's stable preferences from preferences that collapse once their consequences become apparent requires observing that population across a trajectory long enough for the collapse, if it is going to happen, to occur.

This requirement is what finally closes the loop back to Chapter 6's second unmet demand. A civic simulation for geoengineering must expose termination trajectories before adoption, and it can only do so credibly if the simulation itself does not permit costless resets, since a population that has never had to live with the consequences of a failed or withdrawn intervention inside the simulation

has not actually rehearsed the termination-shock scenario Chapter 4 identified as the central risk of any sustained intervention. The simulator's refusal to reset is not a technical limitation to be engineered around. It is the feature that makes its evaluation of termination trustworthy.

What the Simulator Is Not

It remains necessary to state plainly what this architecture does not do, since each of the following clarifications forecloses a misreading this book cannot afford. The civic simulator is not a prediction market; it aims to repair theories by exposing their internal mechanisms, not merely to price competing beliefs against one another. It is not a policy engine; it does not dictate which intervention a population should build, and the normative choice among admissible futures remains, as it must, with the communities who will inhabit the consequences. It is not a game, in the sense of accumulating points toward a terminal victory condition; its objective is the maintenance of admissibility, a condition without a finish line. And it is not a replacement for democracy; it is a consequence-bearing substrate that democratic deliberation can be conducted upon, not an oracle that democratic deliberation can be handed over to.

What it is, correspondingly, is threefold. It is an auditor, producing public, traceable assessments in something like the manner of an environmental impact review, but one that updates as new perturbations are explored rather than freezing at a single filed report. It is a repair engine, actively isolating and resolving the structural disagreements, physical, institutional, model-level, or preference-level, that a purely narrative process would leave permanently tangled. And it is an epistemic environment, maintaining multiple competing models simultaneously without prematurely declaring any one of them correct, which is exactly the discipline Chapter 1 asked for and narrative coordination cannot supply.

The Question Part II Leaves Open

The central claim this chapter inherits from its source framework can now be stated in the vocabulary this book has built across its first two parts. The central question facing any civilization, or any population weighing a planetary intervention, is not which future to choose. It is which futures remain reachable after the choice has been made, evaluated across physical, institutional, and narrative admissibility together, repaired along whichever channel its failures actually occur in, tested against postures it did not design itself around, and persisted through rather than reset away from. A population that cannot pose this question is navigating without a model. A population that can pose it has converted itself, in the framework's own terms, from a collection of declared beliefs into a hypothesis capable of learning from its own consequences.

Parts III through VI now take up four families of intervention in turn, active polar albedo maintenance, oceanic continuation engineering, constructed terrestrial productivity, and planetary material circulation, each of which this chapter's architecture, and Chapter 6's simulation-before-commitment standard, will be applied to directly: not as a rhetorical framing device, but as the explicit test each proposed intervention must be run through before its central claim, that it makes some future reachable, can be trusted rather than merely told.

Part III

Active Polar Albedo Maintenance

Chapter 8

Poles as Control Surfaces

Parts III through VI take up, in turn, the four families of intervention the Introduction promised: polar, oceanic, terrestrial, and planetary-circulation. Each will be run through the explicit procedure Chapter 7 established, resource geometry, admissibility, repair, persistence, and collective choice, and stress-tested against the six governing postures, Optimization, Stability, Exploration, Conservation, Innovation, and Repair. Before that procedure can be applied to a specific polar proposal, however, a prior question needs an answer: why the poles at all. Of every region on Earth where thermal or reflective intervention might in principle be attempted, why should the poles receive the first and most sustained attention this book gives to any single geography.

The answer is not that the poles are cold, nor that they are remote, nor that they are, in the sense Chapter 1 rejected, closer to some optimal planetary temperature than other regions are. The answer is that the poles function, in the coupled Earth system, as control surfaces: small-area, high-leverage regions whose state disproportionately governs the admissible trajectories of systems far larger than themselves, in the same sense that an aircraft's control surfaces, the ailerons, the elevator, the rudder, occupy a small fraction of the airframe's total area yet determine which of the airframe's possible trajectories through the air remain reachable at any given moment. A control surface is not powerful because of its size. It is powerful because of its position within a coupled system whose larger structure amplifies whatever change occurs at that position.

The Ice-Albedo Feedback as an Amplifying Position

The mechanism that gives the poles this position is the ice-albedo feedback, and its structure is worth stating precisely because the interventions in the chapters that follow are all, in different ways, attempts to act directly on it. Sea ice and land ice reflect a large fraction of incident solar radiation back to space; open water and bare ground absorb most of it. When ice retreats, the surface beneath it absorbs more energy, which warms the surrounding region, which drives further retreat. This is a positive feedback in the literal sense, not a beneficial one: a small initial perturbation is amplified rather than damped, and the amplification is strongest exactly where ice coverage is transitioning between its two stable states, present and absent, rather than at points where the region is already fully ice-covered or fully ice-free year-round.

This transitional zone, the region of the ice-albedo feedback's steepest gradient, is where a control-surface intervention has the greatest leverage, because it is where the smallest applied change produces the largest resulting change in the system's trajectory. A given quantity of reflective area

maintained at the margin of a retreating ice sheet, where the feedback loop is actively amplifying small changes, does more work, in terms of the aggregate admissible-volume objective from Chapter 7, than the same quantity of reflective area maintained deep within a region that would remain ice-covered regardless of intervention, or attempted in a region that has already fully transitioned to open water and where the feedback no longer has ice to amplify.

This has an immediate consequence for how the chapters that follow should be read. None of the proposals in this part claims to refreeze the poles, still less the oceans generally; Chapter 10's title states the more modest and more defensible aim directly. The claim is narrower and, this book will argue, correspondingly more admissible: that a targeted intervention at the specific margins where the ice-albedo feedback is actively amplifying loss can arrest or slow that amplification at a cost, and with a footprint, wildly disproportionate to the cost and footprint of any intervention attempting to act on global temperature directly.

Reachability at the Margin

This targeting principle can be restated in the continuation-geometric vocabulary Chapter 4 established. The reachable region of Earth-system trajectories, holding all else equal, is larger when the ice-albedo feedback's amplification is weaker, because a weaker feedback means smaller perturbations propagate into smaller consequences, leaving more of the surrounding state space navigable rather than rapidly foreclosed by runaway loss. A control-surface intervention at the ice margin is therefore, on the objective Chapter 7 formalized as $\text{Vol}_R(A) > 0$, an intervention aimed directly at preserving the volume of admissible futures rather than at moving any single variable, global mean temperature, total ice mass, toward a preferred value. This is the same distinction Chapter 1 drew between a diagnostic variable and an engineering objective, applied now to a specific geography rather than to the planet as an abstraction.

It also clarifies what would count as evidence that a polar intervention is working, as distinct from evidence that some scalar has moved favorably. An intervention that measurably slows the retreat of a specific, monitored ice margin, thereby weakening the local feedback loop at that location, is doing the work this part asks of it, even if global mean sea-ice extent, aggregated across regions the intervention was never targeted at, shows little change. Conversely, an intervention that increases total frozen tonnage somewhere convenient to measure, without affecting the amplifying margins where feedback loss is actually occurring, is producing exactly the kind of maladmitted improvement Chapter 2 warned against in the biomass case: a favorable-looking quantity that does not correspond to the functional relation, in this case a slowed feedback rather than a preserved food web, that actually mattered.

Two Proposals, One Objective, Different Scales

The two proposals evaluated in this part, a nuclear-powered ice-maintenance system taken up in Chapter 9 and a cloud-and-salt hybrid albedo approach taken up in Chapter 10, are not competing solutions to the same problem so much as interventions calibrated to different points along the same control-surface logic. The nuclear-powered system, as Chapter 9 will show, is a high-intensity, spatially concentrated intervention capable of actively maintaining ice at specific, chosen locations

against active thermal pressure, at correspondingly high capital and operational cost and correspondingly deep institutional and jurisdictional exposure. The cloud-and-salt hybrid, as Chapter 10 will show, is a lower-intensity, more distributed intervention aimed not at maintaining ice directly but at reducing the incident radiative forcing across a broader margin, stabilizing what Chapter 10 calls seams, the specific structural junctures, ice-shelf grounding lines, pinning points, narrow channels, whose failure disproportionately accelerates loss in the ice sheets behind them.

Neither proposal is presented as sufficient on its own, and neither will be defended by the single-variable, single-mechanism reasoning Chapter 1 identified as the ordinary and inadequate shape of the geoengineering imagination. Both will instead be run through the same five-step procedure and the same six-posture stress test, so that whatever this part ultimately recommends, if it recommends anything beyond a comparative structure, rests on the same standard the rest of this book applies throughout: not whether a proposal moves a preferred number, but whether it expands the volume of reachable, viable, and recoverable futures available to the systems it touches, and whether its own eventual termination has been planned for with the same seriousness as its operation.

Chapter 9

The Refrigerator-Ice-Machine Proposal

The most technically ambitious polar intervention this book considers is a nuclear-powered ice-maintenance system, informally the refrigerator-ice-machine proposal, that actively removes heat from a targeted ice margin using reactor-driven refrigeration rather than relying on ambient conditions to maintain existing ice or allow new ice to form. Where passive maintenance strategies wait for winter conditions to refreeze what summer has thinned, this proposal treats the margin as a load to be actively managed, in the same sense that a commercial refrigeration system actively manages a cold-storage warehouse regardless of the ambient temperature outside it.

Resource Geometry

The proposal's resource geometry, in the sense Chapter 7's explanatory stack requires as its first step, is dominated by three quantities: the reactor's thermal output available for extraction as usable refrigeration capacity, denoted Φ ; the rate at which that capacity can be transported to and distributed across the targeted ice margin, denoted v ; and the thermal storage the margin itself represents once maintained, denoted S , the buffer of latent heat the ice mass can absorb before its own melting accelerates further loss. A land-based or ice-shelf-anchored reactor installation supplies Φ directly; the distribution problem, getting that refrigeration capacity to bear on an extended, often inaccessible margin rather than only at the reactor's immediate location, is where most of the engineering difficulty and most of the cost concentrates, since a reactor sited at a single point does no useful work on a margin kilometers away unless some transport mechanism, brine circulation, refrigerant piping, or a network of smaller satellite units fed from a central plant, carries that capacity outward.

This immediately exposes a scaling constraint the proposal's advocates have sometimes underweighted. The margin lengths at stake in major ice-sheet and sea-ice retreat are measured in hundreds to thousands of kilometers, while a single reactor installation's economically justifiable distribution radius is measured, at best, in tens of kilometers before transport losses and infrastructure cost dominate the calculation. A single installation, however large, is therefore never a solution to a margin-scale problem; it is, at most, a demonstration of the mechanism at a single control point, and the proposal's actual admissibility depends on whether a network of such installations, sited according to the control-surface logic of Chapter 8, at the specific amplifying junctures rather than uniformly along the margin, can do enough work at enough of those junctures to weaken the aggregate feedback meaningfully.

Admissibility

Physical admissibility, A_P , is comparatively favorable for this proposal: nuclear reactors and industrial refrigeration are both mature technologies, and the thermodynamics of extracting heat from an ice margin and rejecting it elsewhere, most plausibly to the atmosphere or to deep, already-cold water where the rejected heat does comparatively little additional damage, are well understood, though the specific engineering of large-scale brine or refrigerant circulation across an ice-shelf environment, including a sheet that itself moves and fractures, remains a genuinely open problem rather than a solved one.

Institutional admissibility, A_I , is far less favorable, and this is where Chapter 7's stratified-admissibility distinction does its most important work for this proposal specifically. Every major ice margin under active retreat sits within, or adjacent to, a jurisdictional structure poorly suited to authorizing a nuclear-powered installation: the Antarctic Treaty System's restrictions on nuclear material and industrial activity for the southern case, contested and overlapping Arctic sovereignty claims for the northern one. A proposal that is physically executable can therefore remain, for years or decades, institutionally stalled, precisely the gap Chapter 7 warned is routinely misclassified as physical impossibility. Narrative admissibility, A_N , compounds this further: public association of the word nuclear with catastrophic risk, largely independent of this specific proposal's actual risk profile, constitutes a real obstacle to adoption that the proposal's physical soundness does nothing to resolve.

Repair

Of the four repair channels Chapter 7 distinguished, this proposal's likely failure modes concentrate heavily in the institutional and preference channels rather than the infrastructure or model channels. A failure of the reactor itself, a mechanical or thermal fault, is the kind of failure infrastructure repair, in the ordinary engineering sense, is well equipped to address, and nuclear engineering already has mature protocols for exactly this category of repair. A failure to secure cross-jurisdictional authorization, or a collapse of public support following a poorly handled safety incident elsewhere in the nuclear industry, unrelated to this installation but narratively conflated with it, requires institutional repair, rebuilding a regulatory pathway, and preference repair, providing the specific experiences and counterfactuals needed to separate this installation's actual risk profile from the risk profile the public associates with the word nuclear in general. Model repair is relevant chiefly to the network-scaling question raised above: as pilot installations accumulate operating data, the assumption that a given distribution radius and siting density will weaken the targeted feedback by a given amount is exactly the kind of assumption that should be iterated against observed outcomes, with a public, auditable provenance record, rather than defended indefinitely on the strength of its original engineering estimate.

Persistence

This proposal is the clearest case in this part of a genuine termination-shock risk, and Chapter 4's standard applies to it with unusual directness. A margin actively maintained by refrigeration for a period long enough that the broader ice sheet behind it, and any coastal or ecological systems downstream of it, reorganize around its continued presence, cannot simply have that maintenance

withdrawn without producing the delayed collision Chapter 3 described: a physical trend that was gradual, or in this case actively arrested, colliding all at once with a commitment that assumed the arrest was permanent. A reactor installation's decommissioning, whether planned or forced by funding collapse, political rupture, or infrastructure failure, must therefore be modeled, inside the civic simulator Chapter 7 specified, as its own trajectory, with its own recoverability boundary, before the installation is ever built, not discovered afterward as an unplanned second crisis layered on top of whatever climatic trend the installation was meant to arrest.

Collective Choice

The population whose preference must actually be revealed here is not a single national electorate but an overlapping set of jurisdictions, treaty signatories, and downstream-affected communities with no existing forum adequate to the decision's actual shape. This is precisely the condition Chapter 6 argued narrative coordination cannot resolve and Chapter 7's civic simulator is meant to address: a shared, persistent exploration of the proposal's operating and withdrawal trajectories, open to every affected jurisdiction, in which the tradeoff between concentrated high-leverage intervention and concentrated institutional and narrative exposure can be navigated rather than merely argued about across incompatible national narratives.

Stress Test: Six Postures

Run against Chapter 7's six governing postures, this proposal's profile is sharply uneven rather than uniformly strong or weak. **Optimization** favors it strongly, since it accepts the proposal's high upfront cost and institutional friction in exchange for its high per-installation leverage at the amplifying margin. **Conservation** is more ambivalent, favoring the proposal's targeted, small-footprint intervention over a larger-scale thermal approach, but wary of the nuclear material and industrial infrastructure the installation introduces into a previously undeveloped ice environment. **Innovation** favors it, treating current distribution-radius limitations as a problem later engineering iterations, not present-day restraint, should resolve. **Stability** is unfavorable, since the proposal represents a significant deviation from any historical reference trajectory for polar management and introduces a novel dependency the system has no prior experience absorbing. **Exploration** is cautiously favorable, since a successful pilot installation would open a genuinely new region of the reachable-futures space rather than merely reinforcing an already-explored one. **Repair** is the most exposed posture here: this proposal creates, by its own operation, a large new repair burden should it need to be wound down, and a posture prioritizing restoration of admissibility above all other metrics must weigh that created burden heavily against the proposal's operating benefits, arguably more heavily than any of the other five postures are inclined to. This divergence, rather than any single posture's verdict, is the proposal's most important structural feature, and it recurs, in a different configuration, in the alternative Chapter 10 now takes up.

Chapter 10

Holding the Seams: Cloud-and-Salt Hybrid Albedo

Where Chapter 9's proposal actively removes heat from a targeted margin, the intervention this chapter evaluates does not attempt to freeze anything. It attempts, more modestly, to reduce the incident solar radiation reaching a small number of specific structural junctures whose stability disproportionately governs the retreat of the ice sheets behind them, a strategy this book calls holding the seams rather than freezing oceans, to mark clearly that the objective is structural stabilization at a small number of high-leverage points, not thermal reversal across any extended area.

What a Seam Is

A seam, in this sense, is a grounding line, a pinning point, or a narrow ice-shelf channel: a location where an ice sheet's structural integrity depends on contact with, or confinement by, a specific bathymetric or topographic feature, and where loss of that contact or confinement permits an acceleration of flow and calving far beyond what the seam's own small area would suggest. Grounding-line retreat on a reverse-sloped bed is the paradigmatic case: once an ice sheet's grounding line retreats past the crest of an underlying ridge onto ground that deepens inland, the ice sheet is exposed to progressively thicker, more buoyant water as retreat continues, and the retreat becomes self-sustaining, a marine ice-sheet instability, independent of any further external forcing. A seam, in the sense this chapter uses the term, is exactly this kind of instability threshold, and the objective of holding it is to prevent the crossing of the threshold, not to reverse a retreat that has already crossed it.

This makes the seam concept a direct, more specific instance of the control-surface logic Chapter 8 introduced. A seam is where the ice-albedo and ice-dynamical feedbacks are simultaneously steepest, and an intervention that prevents a specific seam's failure does more work, per unit of applied effort, than an intervention distributed uniformly across a margin with no attention to where its structural thresholds actually sit.

The Hybrid Mechanism

The proposed mechanism combines two established but independently insufficient techniques. Marine cloud brightening, seeding low-lying clouds with sea-salt aerosol to increase their droplet number and thereby their reflectivity, can reduce incident radiation over a targeted ocean region during

the specific seasonal windows when solar input is highest and thermal pressure on the seam is most acute, without requiring any permanent installation at the seam itself. Localized surface treatments, reflective materials or engineered snow deployed directly onto the exposed ice near the seam during periods of maximum vulnerability, can further reduce local absorption at the specific structural feature the intervention is targeting, rather than across an undifferentiated area. Neither technique alone reliably prevents a seam's failure; the hybrid approach treats them as complementary, the cloud-brightening component reducing the regional radiative load and the surface component adding a further, spatially precise reduction at the seam itself during the narrow window when it matters most.

Resource Geometry

This proposal's resource geometry is markedly different from Chapter 9's. Its dominant flux, Φ , is not thermal output but aerosol dispersal capacity, the rate at which sea-salt aerosol can be generated and delivered to the target cloud layer, typically from a small fleet of dedicated vessels or platforms rather than a single fixed installation. Its transport rate, v , concerns atmospheric dispersal and cloud-formation timescales rather than brine circulation, and is therefore far more sensitive to weather variability than Chapter 9's proposal, whose refrigeration capacity is largely weather-independent. Its storage term, S , is weakest of the three proposals considered in this book so far: cloud brightening's effect does not accumulate the way frozen mass or processed biomass does, and must instead be sustained through repeated, seasonally timed intervention rather than banked against future need.

Admissibility

Physical admissibility, A_P , for marine cloud brightening at the scales and durations this proposal requires remains genuinely contested rather than settled; the underlying microphysics of aerosol-cloud interaction is well studied in principle but its regional, targeted application at seam-relevant scale has limited field validation, and this uncertainty should be stated plainly rather than assumed away. Institutional admissibility, A_I , is comparatively more favorable than Chapter 9's proposal: no nuclear material is involved, the required vessels operate largely in international or territorial waters with existing precedent for marine research activity, and the intervention's reversibility, in the narrow sense that halting aerosol dispersal halts the radiative effect within days rather than leaving behind a standing installation, reduces the depth of institutional commitment any single authorization represents. Narrative admissibility, A_N , is more favorable for the same reason nuclear framing burdens Chapter 9's proposal, but carries its own distinct risk: any technique broadly classified as solar radiation management inherits public and scientific anxiety about moral hazard, the concern that visible intervention will reduce pressure for emissions reduction elsewhere, an anxiety this book takes seriously without resolving it here, since it belongs properly to Part VII's treatment of governance and authorization.

Repair

This proposal's repair profile inverts Chapter 9's. Infrastructure repair is comparatively cheap and fast, since a damaged or lost dispersal vessel is a far smaller loss than a damaged reactor installa-

tion, and the intervention's lack of permanent structure at the seam itself means there is little fixed infrastructure to repair in the first place. Model repair is comparatively urgent and difficult, precisely because the physical admissibility question above remains open: as deployments accumulate, the gap between predicted and observed regional cloud response must be tracked with the same public, auditable provenance record Chapter 7 specified, since this proposal's central uncertainty sits in exactly this gap rather than in engineering execution. Institutional and preference repair are lighter burdens than for Chapter 9's proposal, given the lower jurisdictional and narrative friction described above, though the moral-hazard concern noted there constitutes its own distinct preference-repair burden, one about the intervention's broader signaling effect rather than about its direct risk.

Persistence

Because this intervention's effect does not accumulate and must be sustained through repeated seasonal deployment rather than banked, its termination-shock profile differs sharply from Chapter 9's. Withdrawal does not confront a system that has reorganized around a standing structure's permanent presence; it confronts, instead, a seam that returns, over the following season, to whatever thermal exposure it would have faced absent the intervention, a return that is itself a form of shock if the interruption is unplanned or comes without warning to the populations and ecosystems whose expectations, over the intervention's operating years, may have adjusted to the seam's continued stability. Persistence, in this case, is less about the standing infrastructure Chapter 4 worried about and more about the continuity of the seasonal commitment itself: an intervention that a population comes to expect every year, and then does not receive, produces its own version of the delayed-collision dynamic Chapter 3 described, even without a single piece of permanent hardware left behind to decommission.

Collective Choice

The population relevant to this proposal's authorization is narrower and more tractable than Chapter 9's, concentrated among the coastal states and research institutions with direct interest in the specific targeted seam, though the moral-hazard concern above means that a wider global public, whose interest lies in whether this class of intervention as a whole reduces pressure for emissions reduction, has a legitimate claim to participate in the civic simulation's exploration even where it lacks direct jurisdiction over the specific seam being targeted.

Stress Test: Six Postures

Against Chapter 7's six postures, this proposal's profile diverges from Chapter 9's in instructive ways. **Conservation** favors it strongly, precisely because of its low permanent footprint and high reversibility relative to Chapter 9's proposal. **Stability** is more favorable here than for Chapter 9, since seasonal, reversible intervention deviates less sharply from historical baseline conditions than a standing nuclear-refrigeration installation does. **Optimization** is more ambivalent than for Chapter 9, since the lower per-deployment leverage and higher scientific uncertainty reduce the confidence with which a favorable cost-benefit case can currently be made. **Innovation** favors it, treating the physical-admissibility uncertainty as exactly the kind of gap further field research should close rather

than a present barrier to deployment. **Exploration** is favorable but more cautious than for Chapter 9, since the intervention's reversibility means a failed exploration is genuinely cheap to abandon, which is itself, on Chapter 6's account of the difference between exploration and commitment, a virtue rather than a limitation. **Repair** is markedly more favorable here than for Chapter 9: because this intervention creates little standing infrastructure and little institutional lock-in, its own repair burden, should it need to be discontinued, is correspondingly small, which is the single clearest structural advantage this proposal holds over the alternative Chapter 9 evaluated.

What Part III Establishes

Taken together, the two proposals evaluated in this part are not rivals to be adjudicated by declaring one superior. They occupy different positions in the tradeoff Chapter 8 anticipated, between concentrated high-leverage intervention with correspondingly deep institutional and repair exposure, and distributed, reversible, lower-leverage intervention with correspondingly shallower exposure but greater scientific uncertainty about whether it works at all. A civic simulation adequate to this choice, run according to Chapter 7's procedure, does not need to resolve which proposal is correct in the abstract. It needs to let the populations actually exposed to a given seam or margin explore both trajectories, including their withdrawal trajectories, and reveal, through that exploration, which tradeoff they are willing to accept for their specific, geographically located stake in the outcome. Part IV now turns from the poles to the ocean more broadly, where a structurally similar tradeoff, between concentrated engineering intervention and distributed biological and material processes, reappears in a different form.

Part IV

Oceanic Continuation Engineering

Chapter 11

Boundary Conditions in Motion

Part III treated the poles as control surfaces: fixed, if retreating, geographic positions where a small applied intervention could disproportionately influence the trajectory of a much larger system. The ocean does not offer this same fixity. It has no analog to a grounding line or a pinning point that stays where it is found; its relevant structures, currents, thermal fronts, storm systems, biomass concentrations, are themselves in continuous motion, and an intervention conceived as acting on a fixed location will frequently discover that the location it needs to act on has already moved by the time the intervention arrives. This part asks what continuation-geometric intervention looks like when the boundary conditions themselves are mobile rather than static, and it does so across four chapters that move from the most literally mobile intervention this book considers, a routed iceberg, through the ocean's own biological productivity, to the processing infrastructure, the Caldera Reactor, that gives this part its name: oceanic continuation engineering, engineering not of a fixed place but of a continuing process.

Why Mobility Changes the Admissibility Calculation

A fixed installation, of the kind Chapter 9 evaluated, accumulates a single, well-defined set of institutional obligations: a jurisdiction it sits within, a permitting body it answers to, a decommissioning plan tied to a specific site. A mobile intervention, an iceberg under tow, a fleet of kelp-cultivation platforms, a reactor vessel that may itself relocate to follow a Sargassum bloom, accumulates a different and in some ways harder problem: it may cross multiple jurisdictions over its operating life, and the admissibility question Chapter 7 formalized as A_I must therefore be evaluated not once, at a fixed site, but continuously, along a route or within a range, as the intervention's position changes relative to the waters it occupies.

This does not make oceanic intervention less admissible than polar intervention in any general sense; it makes the admissibility question differently shaped. A routed iceberg that never enters territorial waters and remains in the high seas for its entire operating life faces a comparatively simple, if still nontrivial, admissibility problem, since the high seas carry their own body of international law rather than the fragmented, jurisdiction-by-jurisdiction problem a coastal deployment would face. A kelp farm anchored within a single nation's exclusive economic zone faces the more conventional, single-jurisdiction problem Chapter 9's reactor installation faced. The Caldera Reactor, evaluated last in this part, sits at both extremes depending on whether it is stationed to process a stable kelp baseline near a fixed farm or dispatched to intercept a mobile Sargassum bloom wherever it happens to be accumulating that season.

Mobility as a Repair Resource

Mobility is not only a complication. It is, on the repair-channel account Chapter 7 developed, itself a resource. A fixed installation that fails catastrophically fails in place, and its infrastructure repair, in the Chapter 7 sense, must occur at that location or not at all. A mobile intervention that begins to fail, an iceberg breaking apart under tow, a cultivation platform damaged by storm, can in principle be relocated, either to a safer holding pattern while repair is arranged or away from a location where its failure would do the most damage, before the failure fully materializes. This is a form of repair unavailable to any of the interventions Part III considered, and it is one reason this book treats the ocean, despite its greater admissibility complexity, as offering genuine advantages over fixed polar or terrestrial intervention for certain classes of problem, particularly problems where the location of greatest need, a specific hurricane's projected path, a specific bloom's current position, is itself not known until shortly before intervention is required.

The Four Chapters Ahead

Chapter 12 takes up the most literally mobile of this part's interventions, routed icebergs used as thermal boundary conditions ahead of forming hurricanes, and confronts directly the unilateralism risk this book's Introduction identified in Stephenson's fictional case. Chapter 13 turns to a form of oceanic productivity that is mobile in a different sense, drifting or cultivated biomass, kelp farmed deliberately and Sargassum accumulating opportunistically, and extends Chapter 2's maladmitted-biomass framework from diagnosis into intervention. Chapter 14 closes the part with the Caldera Reactor, the processing infrastructure that converts both the managed and the opportunistic biomass streams from Chapter 13 into stable, storable, or usable material forms, and is itself capable of the mobility this chapter has argued the ocean rewards. Each chapter applies Chapter 7's full procedure, resource geometry, admissibility, repair, persistence, and collective choice, together with the six-posture stress test, exactly as Part III did, so that the comparison across all six interventions evaluated so far, two polar and three oceanic plus the Caldera Reactor, remains available on consistent terms when Part VII returns to the question of governance across the whole portfolio.

Chapter 12

Icebergs as Mobile Boundary Conditions

Hurricane intensification depends, among other factors, on sea-surface temperature along the storm's projected path: warmer surface water supplies more latent heat to the developing system, and a sufficiently large reduction in surface temperature ahead of a forming or intensifying storm can measurably weaken it. This is the physical premise behind a proposal this book treats with particular care, since it is the proposal that most directly recalls the unilateral installation at the center of Stephenson's *Termination Shock*: towing large icebergs, harvested from calving events already occurring at the margins Part III examined, into the projected path of a forming hurricane, where their melt draws heat from the surrounding surface water and locally lowers the sea-surface temperature the storm would otherwise draw upon.

Resource Geometry

The proposal's flux, Φ , is the cooling capacity available per iceberg, a function of mass and surface area exposed to the surrounding water, and this capacity is consumed, not renewed, over the intervention's operation: an iceberg used to cool one storm's path is, by the time it has done so, substantially melted, and cannot be reused for a second interception without first being replenished from a new calving source. Transport rate, v , concerns towing speed and the lead time available between a storm's formation and its landfall, typically measured in days, which constrains how far a given iceberg can plausibly be moved from its point of origin to a useful interception point within the available window. Storage, S , is essentially the iceberg's own latent heat of fusion, banked at its point of origin and drawn down as it melts in transit and at its destination; unlike Chapter 9's proposal, this storage cannot be replenished in place, and unlike Chapter 10's, it does not need to be, since each intervention is a single-use deployment against a single storm.

This resource geometry immediately exposes the proposal's central practical constraint: the supply of icebergs of useful size, calving at a rate and location that permits towing to a relevant hurricane-formation region within the available lead time, is neither large nor reliably positioned relative to where hurricanes actually form. A proposal that depends on abundant, conveniently located calving events is depending on a supply this book's own Part III has already shown to be a symptom of ice-sheet instability rather than a resource to be relied upon, an uncomfortable tension this chapter does not resolve away.

Admissibility

Physical admissibility, A_P , is reasonably well established for the core mechanism, large-scale iceberg towing has been demonstrated, if not routinely practiced, and the thermodynamics of melt-driven surface cooling are well understood; the open physical question is one of scale, whether a towable iceberg or a small towed fleet can lower sea-surface temperature across an area and by a margin sufficient to measurably weaken a forming storm, rather than being a comparatively minor perturbation against an ocean basin's much larger heat content.

Institutional admissibility, A_I , is where this proposal's kinship with Stephenson's fictional case becomes most direct and most cautionary. A hurricane's projected path crosses, potentially within days, the territorial waters and coastal interests of multiple nations, none of whom may have been consulted before a towing operation begins, and an intervention that measurably weakens a storm approaching one coastline may, depending on the storm's subsequent track, be argued, rightly or wrongly, to have altered its subsequent behavior in ways that affect a different coastline downstream. This is not a hypothetical concern invented for symmetry with the novel; it is the concrete institutional shape any real deployment of this mechanism would immediately confront, and it is exactly the kind of decision this book's Part II argued should never be made unilaterally, by a single well-resourced actor convinced of its own case, but only through the shared, consequence-bearing exploration Chapter 6 and Chapter 7 specified.

Narrative admissibility, A_N , is correspondingly fragile: any weakening intervention that is followed, coincidentally or not, by unexpected damage elsewhere invites an attribution of causation the underlying physics may not support, and a population's willingness to sustain support for the mechanism through such an episode is untested.

Repair

Infrastructure repair is limited in scope here, since the intervention's central asset, the iceberg itself, is consumed rather than damaged in the ordinary engineering sense; there is little to repair once a deployment has occurred beyond assessing towing-fleet readiness for the next one. Institutional repair is the dominant burden this proposal creates, given the multi-jurisdictional exposure just described, and any actual deployment would need pre-negotiated, multilateral authorization protocols in place before a storm forms, not improvised during the narrow window a forming storm allows. Preference repair is closely related: a population's confidence in the mechanism, once shaken by a misattributed downstream event, is not restored by further explanation alone but requires the same kind of sustained, consequence-bearing exploration inside the civic simulator that established the mechanism's case in the first place. Model repair concerns the open scaling question above, and should be treated, at least initially, as unresolved rather than assumed favorable, since a mechanism deployed at real scale before its actual cooling effect on storm intensity has been established against observed outcomes risks the same narrative overreach Chapter 6 warned against.

Persistence

This proposal has an unusually short operational horizon per deployment, days rather than years, which limits the kind of deep, long-accumulating termination shock Chapter 9's standing installation risked. But a different persistence risk applies at the institutional level: if a coastal population comes to expect iceberg interception as a standing service against future storms, on the strength of one or two successful early deployments, the absence of a suitable iceberg at the moment a subsequent storm forms, whether due to supply scarcity or towing-fleet unavailability, produces its own version of the delayed-collision dynamic Chapter 3 described, a population that has organized its expectations, and potentially reduced investment in other forms of storm resilience, around a service that turns out not to be reliably available when needed.

Collective Choice

The relevant population is, again, multi-jurisdictional and only partially overlapping with the population that would benefit most directly from a given deployment; coastal communities along a storm's eventual, post-intervention track have as much stake in the decision as those along its pre-intervention projected path, and a civic simulation adequate to this proposal must give both populations visibility into the same modeled trajectory before any towing operation begins, rather than after.

Stress Test: Six Postures

Optimization is genuinely divided on this proposal in a way it was not for Chapter 9's or Chapter 10's: the intervention's low per-deployment cost relative to its potential damage-avoidance benefit is attractive, but the unresolved scaling question above means the expected benefit itself is poorly characterized. **Conservation** is unfavorable, given the proposal's dependence on a calving supply that Conservation, consistently with its position throughout this book, would prefer to see reduced rather than relied upon. **Innovation** is favorable, treating the scaling and supply questions as engineering problems still open to solution. **Stability** is strongly unfavorable, given the proposal's direct kinship with the unilateral-intervention scenario this book has repeatedly used as a cautionary case. **Exploration** is favorable in principle, since a successful demonstration would open a genuinely new intervention class, but qualifies its endorsement heavily on the multi-jurisdictional authorization problem being solved first, in simulation, rather than in the field. **Repair** is the most skeptical posture here, for reasons distinct from Chapter 9's case: it is not the standing infrastructure that concerns Repair, since there is little of that, but the institutional repair burden a poorly authorized or misattributed deployment would create, a burden Repair judges disproportionate to the mechanism's currently uncertain benefit.

Chapter 13

Kelp Farms and the Sargassum Problem

Chapter 2 established, through the paired krill and Sargassum cases, that biospheric health depends on whether biomass remains distributed, reachable, transformable, and incorporable, not on its aggregate quantity, and set an explicit evaluative standard for kelp cultivation: to be judged by its effects on marine food webs, carbon pathways, oxygen, harvest cycles, and material use, not by growth rate alone. This chapter takes up kelp cultivation and Sargassum interception as a coupled intervention, applying that standard directly, and treats the pairing as the biological half of a system whose processing half, the Caldera Reactor, Chapter 14 will complete.

Resource Geometry

Managed kelp cultivation's flux, Φ , is growth rate per unit cultivated area, a quantity that is comparatively well characterized for established species and cultivation methods, unlike the more volatile and less predictable growth dynamics of an opportunistic Sargassum bloom. Transport, v , concerns harvest logistics, moving cultivated kelp from farm to processing point, a solved problem at modest scale; Sargassum interception's transport problem is harder, since the material must be located, tracked, and intercepted before it reaches shore, a logistics problem closer in kind to Chapter 12's storm-tracking problem than to a conventional agricultural harvest. Storage, S , favors kelp strongly: a managed farm's yield is predictable enough to plan processing capacity around, while Sargassum's storage characteristic is effectively negative, in the sense that the material's ecological cost accrues the longer it sits unprocessed near shore, exactly the incorporation-capacity failure Chapter 2 identified.

Admissibility

Physical admissibility, A_P , is favorable for both streams individually: kelp cultivation is an established practice at meaningful scale, and Sargassum's chemistry and structure, while more variable and occasionally contaminated with heavy metals or plastics as Chapter 2 noted, are not fundamentally intractable to processing. The harder admissibility question is coupling the two streams into a single processing infrastructure without one distorting the incentives governing the other, precisely the concern the memory of this project's own development has already flagged: that a processing system economically dependent on Sargassum's continued anomalous abundance would create an institutional interest in the persistence of the ecological failure it was meant to relieve. This book's position, consistent with that earlier concern, is that kelp must remain the baseline feedstock and

Sargassum a secondary, disturbance-responsive stream, and this allocation is itself part of what any admissibility evaluation of the coupled system must verify, not an incidental design preference.

Institutional admissibility, A_I , favors kelp cultivation, which occurs within established aquaculture regulatory frameworks in most coastal jurisdictions, over Sargassum interception, which, depending on where interception occurs, may cross into international waters and require coordination among multiple coastal states along the bloom's drift path, echoing Chapter 12's multi-jurisdictional problem though at a lower intensity, since interception vessels create no comparable diplomatic exposure to a towed iceberg or a targeted hurricane intervention. Narrative admissibility, A_N , is favorable for kelp, an already broadly accepted practice, and more mixed for Sargassum interception, where public understanding of the material's dual character, valuable offshore habitat, damaging onshore waste, remains limited enough that the intervention risks being narratively flattened into either straightforward cleanup or unnecessary interference, neither of which captures Chapter 2's actual argument.

Repair

Infrastructure repair for kelp farms follows established aquaculture practice and is comparatively routine. For Sargassum interception, infrastructure repair is complicated by the fleet's need to operate wherever the bloom currently is, which may be far from a fixed maintenance base, a mobility cost this part's opening chapter identified as a general feature of oceanic intervention. Institutional repair is most needed where interception crosses jurisdictions without pre-existing coordination agreements. Model repair is the central ongoing task for the Sargassum stream specifically, since bloom formation drivers remain, as Chapter 2 noted, only partially understood, and interception planning that assumes a bloom's future drift path without continuously revising that assumption against observed drift is repeating exactly the single-mechanism overconfidence Chapter 1 warned against. Preference repair matters chiefly for correcting the narrative flattening just described, ensuring affected coastal populations understand the intervention as targeted incorporation-capacity relief rather than either full remediation or unwarranted intrusion.

Persistence

Kelp cultivation's persistence risk is low and conventional: an established farming operation that ceases returns cultivated area to an unmanaged state without the kind of reorganization-around-permanence risk Chapter 9's installation created. Sargassum interception's persistence risk is more particular: coastal communities that come to rely on interception to prevent beach-fouling and near-shore hypoxia may reduce their own local mitigation capacity, cleanup crews, monitoring programs, in the expectation that interception will continue indefinitely, producing exactly the commitment-mismatch structure Chapter 3 described if interception capacity is later reduced or redirected, for instance toward a differently located bloom in a given season.

Collective Choice

Kelp cultivation's relevant population is local and largely conventional, the farm's licensing jurisdiction and immediate coastal stakeholders. Sargassum interception's relevant population is regional

and shifting from year to year with the bloom's location, requiring the same kind of standing, pre-negotiated multi-jurisdictional framework Chapter 12 argued was necessary for iceberg deployment, though at considerably lower diplomatic stakes.

Stress Test: Six Postures

Conservation strongly favors kelp cultivation, an established low-footprint practice, and is more cautious about Sargassum interception's fleet-based footprint, though it recognizes interception as preventing a larger ecological cost than it creates. **Optimization** favors both streams, given their relatively low per-unit cost of intervention relative to the value of the processed output and, for Sargassum, the avoided nearshore damage. **Stability** favors kelp strongly and is neutral to mildly unfavorable toward Sargassum interception, an intervention responding to a phenomenon, anomalous bloom abundance, that is itself a deviation from historical baseline. **Innovation** favors both, particularly the processing technology Chapter 14 will detail. **Exploration** is most interested in the Sargassum stream specifically, since it represents the less-understood and more scientifically informative of the two feedstocks. **Repair** favors the coupled system's overall design, precisely because the kelp-baseline, Sargassum-secondary allocation this chapter defended is itself a repair-oriented design choice, structured to avoid creating a standing dependency on ecological pathology, the specific failure mode Repair is most attentive to across every intervention this book has evaluated.

Chapter 14

The Caldera Reactor: Thermopneumatic Transformation of Maladmitted Biomass

The kelp and Sargassum streams Chapter 13 evaluated require a processing infrastructure capable of converting wet, variable, and in the Sargassum case sometimes contaminated marine biomass into stable, storable, or directly usable material and energy products. This chapter evaluates that infrastructure, an oceanic thermopneumatic system this book calls the Caldera Reactor, whose function is precisely the transformation Chapter 2 argued was necessary wherever biomass exists in a quantity or location a receiving system cannot incorporate: not disposal, and not indefinite storage, but conversion into a different, admissible material form.

The Mechanism

The reactor's core mechanism is a geothermal-steam-driven lift of a large, approximately twelve-meter titanium-ceramic compression plate, whose rise draws seawater upward through condensation-driven vacuum, carrying with it layered kelp, peat, and sediment feedstock into a compression chamber beneath the plate. The plate's descent compresses this layered material under substantial pressure, with the compression stroke's mechanical energy partially recovered through an integrated turbine rather than dissipated as waste heat. Biomass entering the system is classified in real time by a Raman-sensing and convolutional-neural-network pipeline, distinguishing clean, cultivated kelp from degraded or contaminated Sargassum loads and routing each toward the processing pathway suited to it, hydrous pyrolysis for high-moisture feedstock such as kelp, avoiding the energy cost of pre-drying, and anhydrous pyrolysis via a separate thermal pathway for drier or dewatered material, since the two processes exploit genuinely different reaction kinetics, sub- and supercritical water phase behavior in the hydrous case, and cannot be treated as a single process with a variable moisture parameter. Fine-grained control across the reactor's many parallel processing pathways is distributed through a pressure-sensitive thermal-clutch knot lattice, a decentralized ternary fluidic control scheme rather than a single centralized controller, and a biological conversion stage employing an engineered yeast strain, ARX-X27, further processes intermediate output into biocrude, biochemical feedstocks, and polyhydroxyalkanoate bioplastics.

Resource Geometry

The reactor's flux, Φ , is throughput capacity, the volume of feedstock the compression-and-pyrolysis pathway can process per cycle, bounded jointly by the mechanical cycle time of the plate's rise and fall and by the classification pipeline's ability to sort incoming material fast enough to keep both pyrolysis pathways supplied. Transport, v , concerns feedstock delivery, favorable and predictable for kelp delivered from a nearby managed farm, unfavorable and variable for Sargassum delivered from wherever Chapter 13's interception fleet has most recently located a bloom, which is precisely why this chapter's opening framing part argued the reactor should itself be capable of the mobility Chapter 11 identified as an oceanic advantage, relocating toward a stable kelp-farm mooring for baseline operation but redeployable toward an active bloom when Sargassum surplus demands it. Storage, S , is favorable for the reactor's output products, biocrude and bioplastics are both stable, storable forms, in sharp contrast to the near-zero effective storage of the untreated feedstock Chapter 13 described.

Admissibility

Physical admissibility, A_P , rests on established pyrolysis and biological-conversion chemistry, though the specific integration, the thermopneumatic lift-and-compression mechanism, the dual hydrous-anhydrous pathway split, the knot-lattice control scheme, is a novel engineering synthesis whose components are individually plausible but whose combined performance at full scale has not been demonstrated, and this book states that limitation directly rather than presenting the reactor as a solved design. Institutional admissibility, A_I , depends heavily on the reactor's operating mode: a fixed installation moored near a managed kelp farm sits within a single jurisdiction's aquaculture and industrial-processing regulatory frameworks, comparatively tractable; a mobile deployment intercepting an active Sargassum bloom inherits the multi-jurisdictional complexity Chapter 13 already identified for the interception fleet itself. Narrative admissibility, A_N , is favorable insofar as the reactor is legible as waste-to-value conversion, a broadly accepted category of technology, but requires the same care Chapter 13 flagged against being narratively flattened into either simple recycling, understating the reactor's role in preventing the specific incorporation-capacity failures Chapter 2 diagnosed, or an untested speculative machine, understating the individual maturity of its component technologies.

Repair

Infrastructure repair is the reactor's most demanding channel, given the mechanical complexity of the lift-compression cycle and the number of parallel processing pathways the knot-lattice control scheme must coordinate; a fault in any single pathway should, by the decentralized design's own logic, be isolable and repairable without halting the entire system, and this isolability should itself be treated as a design requirement subject to verification, not an assumed property of decentralization in general. Model repair concerns primarily the classification pipeline, whose Raman-sensing and CNN components must be continuously validated against observed contamination and degradation patterns in incoming Sargassum loads, since a misclassification that routes contaminated material into the wrong pyrolysis pathway, or into the biological conversion stage where the ARX-X27 yeast

strain may be more sensitive to contaminants than the thermal pathways are, represents a failure mode with consequences disproportionate to its likely frequency. Institutional and preference repair follow the same channels Chapter 13 described for the interception fleet the reactor depends on for its Sargassum stream.

Persistence

A reactor installation moored near a managed kelp farm for an extended operating period creates a persistence risk structurally similar to, though smaller in scale than, Chapter 9's polar installation: the farm's own harvest planning, and potentially onshore product-distribution infrastructure built around the reactor's steady output, would need to reorganize around the reactor's continued operation, and its decommissioning or relocation should be planned, inside the civic simulator, before the dependent onshore infrastructure is committed to, not after. A mobile deployment intercepting Sargassum blooms carries a different and generally milder persistence risk, since, as Chapter 13 noted, the intervention itself is disturbance-responsive rather than standing, and a population's expectations are correspondingly less likely to organize around its permanent presence at any single location.

Collective Choice

The relevant population divides, as the admissibility discussion above already implied, between the kelp farm's local jurisdiction, whose stake resembles Chapter 13's kelp-cultivation case, and the broader, shifting set of coastal jurisdictions along whatever Sargassum drift path the reactor is currently servicing, whose stake resembles Chapter 13's interception case. A civic simulation exploring the reactor's adoption should model both operating modes and the transition between them, since the reactor's actual value proposition depends on its capacity to serve both, not on either mode considered alone.

Stress Test: Six Postures

Innovation favors the reactor most strongly of any intervention evaluated in this book so far, precisely because of its novel engineering synthesis and clear pathway to iterative improvement as operating data accumulates. **Optimization** is favorable, given the reactor's conversion of an otherwise wasted or damaging biomass stream into stable, valuable products. **Conservation** is cautiously favorable, endorsing the reactor's role in preventing incorporation-capacity failures while remaining watchful, exactly as this chapter has insisted throughout, that the reactor not become economically dependent on Sargassum bloom persistence. **Exploration** favors the reactor strongly, since it opens both a new processing technology and, through its classification and control systems, a new source of data about Sargassum's own variability. **Stability** is the most cautious posture, given the design's genuine novelty and the absence of a directly comparable prior installation to draw confidence from. **Repair** is favorably disposed on balance, given the design's built-in isolability and the deliberate kelp-baseline allocation, but withholds full endorsement until the isolability claim above has been verified rather than merely asserted, a qualification consistent with Repair's posture throughout this book: readiness to trust a design only after its capacity for repair has itself been demonstrated, not assumed.

What Part IV Establishes

Across Chapters 12 through 14, oceanic continuation engineering has proven structurally different from the polar interventions of Part III in a specific and recurring way: every oceanic intervention this part evaluated gains something from mobility, an iceberg that can be towed, a fleet that can track a bloom, a reactor that can relocate between a stable farm and an active bloom, and every one of them pays for that mobility in a correspondingly more complex, continuously renegotiated admissibility problem rather than the single, fixed-site problem Part III's polar installations faced. Part V now turns to a domain where mobility is largely absent and the opposite tradeoff applies: constructed terrestrial productivity, where an intervention's fixed location is not a complication to be managed but the entire basis of what it can accomplish.

Part V

Constructed Terrestrial Productivity

Chapter 15

Substrate Before Growth

Part IV closed by noting that oceanic intervention gains leverage from mobility and pays for it in continuously renegotiated admissibility. Terrestrial intervention, as this part will show, largely inverts that trade. A rainforest does not migrate to follow a favorable microclimate the way a Sargassum bloom drifts with a current, and a support structure built to expand a forest's physical state space is, almost by definition, fixed to the ground it is built on. What terrestrial intervention loses in mobility, however, it gains in a different resource this book has not yet had occasion to emphasize: durability of substrate. A trellis, an anchor, a biochar-lined berm, once built, continues doing its work for decades without needing to be relocated, refueled, or redeployed, and the admissibility question it raises is asked once, at construction, rather than continuously, as Part IV's mobile interventions required.

What a Rainforest Generator Is, and Is Not

The central intervention this part evaluates, in Chapter 16, is a class of support architecture this book calls rainforest generators: tensile trellises, hanging cable grids, smart anchors, engineered light wells, misting systems, and nutrient-fog delivery, deployed not to manufacture a forest from nothing but to expand the physical state space within which existing rainforest growth processes can occur. This distinction matters enough to state plainly before the chapter's technical detail: a rainforest generator does not grow trees. It removes specific physical constraints, insufficient vertical support for climbing and epiphytic growth, uneven light penetration to lower canopy layers, moisture stress during dry-season intervals, that would otherwise cap how much of a given area's growth potential is actually realized. The generator is closer, in the vocabulary this book has used throughout, to widening a viable region than to manufacturing a state directly within it.

Terra Preta Rain

Chapter 17 pairs the generator architecture with a nutrient and moisture delivery system this book calls terra preta rain: rainwater collected in elevated basins, including the calderas this book's oceanic material has already made structurally familiar, filtered through gravel-biochar-sand media and biochar-lined berms, enriched with nutrients, minerals, microbes, and suspended carbon, and redistributed by gravity over the constructed canopy as a continuously circulating medium rather than a static reservoir. The name deliberately invokes the historical Amazonian terra preta soils, anthropogenic dark earths enriched with charcoal and organic matter that have sustained elevated fertility for centuries after their original cultivation ceased, as a precedent for durable, biochar-mediated soil

enrichment, while the "rain" in the name marks the departure from that precedent: a continuously circulating delivery system rather than a one-time soil amendment.

Why This Part Follows Chapter 2's Standard Directly

Chapter 2 was explicit that rainforest generators would be evaluated by habitat complexity, nutrient continuity, and increasing ecological autonomy, not by vertical biomass accumulated. This part inherits that standard without modification, and it matters more here than anywhere else in the book so far, because a support architecture whose stated purpose is to expand growth is exactly the kind of intervention most tempted to report its own success in tonnage, total biomass added, when the standard this book has defended since Chapter 2 requires asking a harder question: has the added biomass become an autonomous, self-sustaining part of the forest's own food web and nutrient cycle, or does it remain a standing dependency on the generator's continued mechanical operation, a distinction with direct consequences for the persistence analysis Chapter 16 will need to perform.

Two Chapters, One Coupled System

Chapter 16 evaluates the trellis-and-anchor generator architecture on its own terms, applying Chapter 7's full procedure. Chapter 17 evaluates terra preta rain as its complementary delivery system, and closes by returning explicitly to the coupling this book's own development notes have already identified: the generator and the delivery system as one half of a larger material loop whose other half, the Caldera Reactor's marine biomass conversion, was evaluated in Chapter 14. That coupling, converted marine carbon routed into terrestrial nutrient delivery, is not a minor cross-reference. It is the clearest instance in this book of Chapter 7's reachability-volume objective operating across, rather than within, a single intervention: neither the reactor nor the rainforest generator alone expands the reachable-futures volume by as much as the two systems coupled together do, and Chapter 17's closing section makes that coupling explicit rather than leaving it implicit across two separately evaluated chapters.

Chapter 16

Rainforest Generators: Expanding the Growth State Space

Resource Geometry

The generator's flux, Φ , is structural load-bearing capacity, the tensile trellis and cable grid's ability to support climbing and epiphytic biomass that unsupported ground-level growth could not sustain, together with light-delivery capacity through the engineered light wells that redirect canopy-level sunlight to strata that would otherwise remain shaded. Transport, v , is comparatively minor for this intervention relative to every prior chapter in this book, since a fixed structure does not move material across space in the way a distribution network or a towed iceberg does; the closer analog to a transport rate here is the rate at which the misting and nutrient-fog systems can deliver moisture and dissolved nutrients across the structure's covered area, a rate bounded by pumping capacity and delivery-line density rather than by any external logistics constraint. Storage, S , is the structural system's own standing capacity, unlike every prior chapter's storage term, this one does not deplete through use; the trellis continues bearing load and the light wells continue redirecting sunlight for as long as the structure itself remains intact, making this intervention's storage term closer to a fixed capital asset than to a consumable resource.

Admissibility

Physical admissibility, A_P , is favorable in the narrow engineering sense, tensile support structures, engineered light management, and misting systems are all individually mature technologies, and the genuine open question is not whether the components work but whether their combined deployment at forest scale, across the uneven terrain and existing canopy a real rainforest presents, can be achieved without the installation process itself causing more disturbance to existing growth than the eventual generator's benefit repays. This is a distinct and more tractable uncertainty than Chapter 10's open physical-admissibility question, since it concerns installation methodology rather than the underlying mechanism's basic viability.

Institutional admissibility, A_I , depends heavily on jurisdiction and land tenure in a way none of this book's prior interventions have confronted directly: a rainforest generator's site is, unlike an ice margin or the high seas, very likely to fall within a specific nation's sovereign territory, frequently territory with existing indigenous or local community claims, conservation designations, or

extractive-industry pressure already contesting its use. This book takes the position, consistent with the placement-objective work of Chapter 5, that a generator sited without the consent and continued participation of the communities whose land and forest it occupies is not merely institutionally fragile but is failing the equity and jurisdictional-distance terms Chapter 5's own objective function made explicit; a generator imposed on a community rather than built with one inherits every risk Chapter 5 warned against when a distribution facility's siting encodes an indefensible prior judgment under cover of a technically optimized location.

Narrative admissibility, A_N , is genuinely double-edged. A visible, engineered structure inserted into a rainforest canopy is legible, favorably, as active restoration, but is equally legible, unfavorably, as artificial intrusion into a natural system, a framing this book's own insistence that the generator expands rather than manufactures growth is specifically intended to correct, though correcting a narrative framing is, as Chapter 6 argued throughout, a matter of sustained exploration rather than assertion.

Repair

Infrastructure repair is this intervention's most conventional channel: a damaged trellis section, a fouled misting line, a failed smart anchor are engineering-maintenance problems with established solutions, complicated mainly by the difficulty of accessing repair sites within a mature forest canopy rather than by any novelty in the repair task itself. Institutional repair concerns primarily the land-tenure and consent question above: a generator whose community relationship deteriorates after construction, whether through inadequate initial consultation or a later change in local governance, requires institutional repair of exactly the kind Chapter 7 specified, rebuilding a cooperative framework, not merely fixing a physical fault. Model repair is centered on the ecological-autonomy question Chapter 15 raised: distinguishing biomass that has become self-sustaining from biomass still dependent on the generator's mechanical operation requires ongoing observation and revision of whatever model initially predicted the transition, since an overconfident early claim of autonomy, unrevised against actual observed dependency, would be exactly the kind of unexamined success-reporting Chapter 2's standard was designed to prevent. Preference repair addresses the narrative double-edge described above, providing the sustained, observable evidence needed for a skeptical local or global public to update away from an intrusion framing without simply being told to.

Persistence

This is the sharpest persistence question this book has posed since Chapter 9, and it is sharper here because it is graded rather than binary: the more of the generator-supported biomass that becomes genuinely autonomous, integrated into the forest's own nutrient cycling and no longer dependent on misting or nutrient-fog delivery, the smaller the termination shock the generator's eventual removal or failure would create, while biomass still dependent on mechanical delivery at the time of removal would face a collapse resembling, in miniature, the marine ice-sheet instability Chapter 10 described: a support withdrawn from a system that has come to depend on it, faster than the system can adapt. This grading is precisely why Chapter 15 insisted the ecological-autonomy standard could not be treated as a one-time success claim; it must be tracked continuously, as a fraction of supported

biomass that has transitioned to autonomy, for the persistence analysis to mean anything at the point termination is actually considered.

Collective Choice

The relevant population is the most geographically concentrated and, in principle, the most directly consultable of any intervention this book has evaluated: the specific community or communities whose land the generator occupies, together with whatever broader conservation or scientific interest a particular forest's ecological significance attracts. A civic simulation exploring this intervention's adoption should give the local population primary standing in the exploration, consistent with the consent principle already argued for above, while making the trajectory's outputs visible to the broader interested public whose scientific or conservation stake is real but secondary to the local community's direct stake in the land itself.

Stress Test: Six Postures

Conservation is favorably disposed in principle, given the intervention's stated aim of expanding rather than replacing natural growth processes, but withholds full endorsement pending resolution of the installation-disturbance question raised under admissibility above. **Exploration** favors the intervention strongly, since it opens a genuinely new class of ecological support architecture with wide potential application across degraded or marginal forest sites. **Innovation** is favorable for the same reason, treating the installation-disturbance and autonomy-tracking questions as problems suited to iterative engineering refinement. **Stability** is more favorable here than for almost any prior intervention in this book, precisely because the generator's aim is to support an existing, historically continuous ecological process rather than to introduce a novel dynamic, a rare case of Stability and Exploration substantially agreeing. **Optimization** is favorable, given the intervention's comparatively low ongoing operating cost relative to a fixed structure's long service life. **Repair** is the most consequential posture for this intervention specifically, since its central concern, the fraction of supported biomass that has genuinely transitioned to autonomy, is exactly the metric this chapter's persistence analysis identified as decisive, and Repair's insistence on tracking that metric continuously, rather than accepting an early or aggregate claim of success, should be read as this chapter's central methodological requirement rather than merely one posture's opinion among six.

Chapter 17

Terra Preta Rain and the Marine-Terrestrial Loop

The Delivery System

Terra preta rain supplies the rainforest generator evaluated in Chapter 16 with the nutrient and moisture input its misting and fog-delivery systems distribute. Rainwater is collected in elevated basins, positioned to exploit existing topography including, where available, volcanic caldera formations whose natural catchment geometry reduces the engineering burden of collection relative to a purpose-built basin. Collected water passes through a filtration sequence of gravel, biochar, and sand layers, and through biochar-lined berms positioned along the water's path to the generator's delivery network, accumulating dissolved nutrients, minerals, beneficial microbial populations, and suspended carbon as it passes. The resulting enriched water is gravity-fed to the generator's misting and fog systems, where it is distributed across the supported canopy as a continuously circulating medium, drawn again from whatever fraction returns to the collection basins through throughfall and runoff, rather than as a single applied treatment exhausted after one use.

Resource Geometry

The system's flux, Φ , is filtered-water throughput, bounded by basin catchment area and by the hydraulic capacity of the gravel-biochar-sand filtration sequence, which must be sized to avoid becoming the bottleneck that limits how much of the generator's delivery capacity can actually be supplied. Transport, v , is gravity-driven distribution rate from basin to canopy, favorable wherever the collection basin sits at sufficient elevation above the generator's delivery network, and this elevation dependency is why the system's siting is coupled tightly to the generator's own siting rather than independently optimizable. Storage, S , is the biochar media's own nutrient- and carbon-retention capacity, which, unlike most storage terms in this book, improves rather than depletes with sustained use over an initial period, as biochar's porous structure accumulates a stable microbial and mineral loading that increases its filtration effectiveness before eventually requiring renewal.

Admissibility

Physical admissibility, A_P , rests on established biochar filtration and soil-amendment chemistry, extended here to a continuously circulating rather than single-application context; the historical terra preta precedent provides genuine evidence for biochar-mediated fertility's long-term durability, though that precedent concerns static soil amendment rather than continuous water-mediated circulation, and this chapter states that extension as a hypothesis the system's own operating data must confirm rather than as an established transfer of the precedent's evidentiary weight. Institutional and narrative admissibility for the delivery system largely inherit Chapter 16's analysis, since the two systems share a site and a community relationship, though the water-collection basin, particularly where it exploits a caldera formation, may introduce an additional land-use or water-rights question distinct from the generator's own footprint, one this book flags rather than resolves, since water rights regimes vary sufficiently by jurisdiction that no general treatment would serve every actual site.

Repair

Infrastructure repair concerns the filtration sequence and delivery plumbing, comparatively conventional maintenance tasks. Model repair is the more interesting channel here: the assumption that biochar filtration effectiveness improves with sustained use before eventually requiring renewal is itself a claim that should be tracked against observed filtration performance over the system's operating life, since biochar media saturated with contaminants rather than beneficial loading would degrade rather than improve, and distinguishing these two trajectories early is exactly the kind of auditable model revision Chapter 7 specified. Institutional and preference repair follow Chapter 16's analysis directly, given the shared site and community relationship.

Persistence

The delivery system's persistence risk is closely coupled to Chapter 16's generator: because the generator's supported biomass depends on continued nutrient and moisture delivery until it reaches the autonomy threshold Chapter 16 tracked, an interruption to the delivery system before that threshold is reached compounds directly into the generator's own termination-shock risk, rather than constituting an independent risk of its own. This coupling is a further reason, alongside the elevation-dependent siting noted above, to treat the generator and the delivery system as a single intervention for persistence-analysis purposes rather than as two independently evaluable ones.

Collective Choice

Identical to Chapter 16's analysis, given the shared site and population; this chapter does not introduce a separate collective-choice question.

Stress Test: Six Postures

The delivery system's stress-test profile tracks Chapter 16's closely across all six postures, given the tight coupling just described, with one distinction worth noting: **Innovation** is more strongly favorable toward the delivery system specifically than toward the generator structure, since the continuous-circulation extension of the terra preta precedent is the more genuinely novel scientific claim in the coupled system, and the one whose confirmation would carry the broadest applicability beyond this specific intervention, to soil and water management contexts far removed from rainforest support architecture.

Closing the Marine-Terrestrial Loop

This chapter closes Part V by making explicit a coupling this book's own development has treated as a standing design intention since before Part IV was drafted: the Caldera Reactor of Chapter 14 converts wet marine biomass, managed kelp as baseline feedstock and intercepted Sargassum as a disturbance-responsive secondary stream, into stable carbon and nutrient forms, biocrude, biochemical feedstocks, and PHA bioplastics among them, some fraction of which can be routed, as processed carbon and nutrient input, into the terra preta rain system this chapter has just evaluated, entering the collection basin's filtration sequence alongside the biochar, minerals, and microbial content already specified there.

This coupling is the clearest instance in this book of two intervention families, evaluated across two separate parts and four separate chapters, jointly expanding the reachable-futures volume Chapter 7 identified as the book's governing objective by more than either could alone. The reactor, without a terrestrial destination for its processed carbon output, converts a marine liability into a storable asset with no specified use beyond storage or conventional sale; the rainforest generator and its delivery system, without a marine-processed carbon input, must rely entirely on whatever the local basin's own catchment supplies, a more limited nutrient budget than the coupled system provides. Neither deficiency is severe enough to make either intervention non-viable alone, and this book does not claim the coupling is necessary. It claims only that the coupling is available, that it was a design intention from early in this project's own development, and that a civic simulation evaluating either intervention in isolation, without exposing the coupled trajectory as an alternative for comparison, would be withholding from the populations doing the evaluating exactly the kind of cross-system information Chapter 7's admissibility framework was built to surface rather than obscure.

What Part V Establishes

Terrestrial intervention, across both chapters of this part, has proven to trade Part IV's mobility-versus-admissibility tension for a different one: fixed, durable structure against a graded, continuously tracked persistence risk tied not to standing infrastructure alone, as in Chapter 9's polar case, but to the fraction of supported biological process that has or has not yet become genuinely autonomous. This graded persistence standard, together with the marine-terrestrial coupling just established, gives Part VI, which now returns to finish the planetary material-circulation infrastructure begun in Chapter 5, a concrete instance of exactly the kind of coupled, cross-part intervention portfolio that infrastructure will need to route, monitor, and govern at planetary scale.

Part VI

Planetary Material Circulation

Chapter 5

Placement Without Disposal

Every account of planetary material circulation eventually confronts a question that sounds narrowly logistical but is not: where should recoverable matter be held so that it remains available for use? The question is usually asked in its negative form, where should a society put what it no longer wants, and answered accordingly, by siting disposal facilities as far as possible from the populations that generate the material being disposed of. This chapter asks the question in its positive form instead. In a civilization organized around repair, reuse, and moderate throughput rather than disposal and replacement, the central logistical institution is not the dump but the distribution reserve: a facility, or a network of facilities, where components, tools, food, medicines, construction materials, and recoverable objects are held near enough to demand to remain useful, but concentrated enough to be inspected, maintained, and redistributed efficiently. The chapter's aim is to show that the siting of such reserves is not a matter of finding cheap land, but a decision that encodes, whether or not it is stated explicitly, a set of judgments about whose need matters, how much redundancy is prudent, and what a society is willing to trade against transport efficiency.

Why Placement Is Not a Fixed Problem

The difficulty of siting a distribution reserve changes with population, but not in a way that reduces to a single scaling law. A civilization of two billion people can tolerate longer transport distances, larger ecological buffers around any given facility, and a comparatively sparse network of reserves, because the consequences of any single placement error are diluted across a smaller and more loosely coupled population. A civilization of ten billion requires much finer placement, because small errors in distance, redundancy, spoilage, or unequal access are multiplied across a population an order of magnitude larger, and because the margin for improvisation around a poorly placed facility shrinks as demand density rises.

Population size alone does not determine the answer, however, and treating it as though it did would repeat the error this book has already identified in the temperature and biomass cases: the substitution of a single scalar for a genuinely multidimensional problem. Density, transport infrastructure, climate, regional repair capacity, political and jurisdictional boundaries, exposure to compound disasters, the physical durability of stored materials, and the spatial distribution of need all bear on where a reserve belongs, and they do not trade off against one another in any fixed proportion. The facility with the largest catchment should not automatically be sited where land is cheapest, nor should it be sited primarily to keep unwanted material at the greatest possible distance from the most affluent populations, since both criteria optimize a proxy (land cost, or distance from a privileged

population) rather than the quantity that actually matters: the total distance, summed appropriately across every affected community, between material, need, repair capacity, and eventual return to circulation.

A Formal Model of Admissible Placement

The judgment involved can be made precise, though making it precise is itself part of the argument, since the resulting expression will turn out to be irreducibly long. Let $i \in I$ index communities, $j \in J$ index candidate distribution sites, $k \in K$ index classes of goods, $t \in T$ index time periods, and $\omega \in \Omega$ index plausible disruption scenarios, each occurring with probability p_ω . Let $x_j \in \{0, 1\}$ indicate whether a distribution reserve is established at site j ; y_{ijkt}^ω the quantity of good k sent from site j to community i in period t under scenario ω ; r_{jkt}^ω the quantity of good k repaired or remanufactured at site j ; q_{ijkt}^ω the quantity returned from community i to site j for repair or reclamation; s_{jkt}^ω the inventory of good k retained at site j ; and u_{ikt}^ω the unmet need for good k in community i .

A socially admissible placement is then the solution to a deliberately expanded objective, one that resists collapsing transport cost, equity, resilience, and ecological exposure into a single term:

$$\begin{aligned}
(x^*, y^*, q^*, r^*, s^*, u^*) \in \arg \min_{x, y, q, r, s, u} \sum_{\omega \in \Omega} p_{\omega} & \left[\alpha \sum_{t, i, j, k} d_{ij} m_k (y_{ijkt}^{\omega} + q_{ijkt}^{\omega}) \right. \\
& + \beta \sum_{t, i, j, k} e_{ijkt}^{\omega} (y_{ijkt}^{\omega} + q_{ijkt}^{\omega}) \\
& + \gamma \sum_{t, i, k} \nu_{ik} u_{ikt}^{\omega} \\
& + \delta \sum_{t, k} \left[\max_{i \in I} \left(\frac{\sum_j y_{ijkt}^{\omega}}{D_{ikt}^{\omega} + \varepsilon} \right) - \min_{i \in I} \left(\frac{\sum_j y_{ijkt}^{\omega}}{D_{ikt}^{\omega} + \varepsilon} \right) \right] \\
& + \eta \sum_{t, i, k} P_i \left| \tau_{ikt}^{\omega} - \frac{\sum_{i'} P_{i'} \tau_{i'kt}^{\omega}}{\sum_{i'} P_{i'}} \right| \\
& + \theta \sum_{t, j} \left[\max \left(0, \sum_{i, k} y_{ijkt}^{\omega} - C_{jt}^{\omega} x_j \right) \right]^2 \\
& + \kappa \sum_{t, j, k} \left(\frac{s_{jkt}^{\omega}}{\sum_{j'} s_{j'kt}^{\omega} + \varepsilon} \right)^2 \\
& + \lambda \sum_{t, i, k} \max \left(0, \rho_{ik} - \sum_j a_{ij}^{\omega} \frac{y_{ijkt}^{\omega}}{D_{ikt}^{\omega} + \varepsilon} \right) \\
& + \mu \sum_{t, j, k} [w_k s_{jkt}^{\omega} + \chi_k \max(0, s_{jkt}^{\omega} - \bar{s}_{jk})] \\
& + \xi \sum_{t, j} h_j \left(x_j + \frac{\sum_{i, k} y_{ijkt}^{\omega}}{C_{jt}^{\omega} + \varepsilon} \right) \\
& + \pi \sum_{t, j} \left(b_j x_j + \zeta_j \sum_{i, k} y_{ijkt}^{\omega} \right) \\
& + \sigma \sum_{t, j, k} \max(0, r_{jkt}^{\omega} - R_{jkt}^{\omega}) \\
& + \phi \sum_{t, i, j, k} y_{ijkt}^{\omega} (1 - \ell_{ijkt}^{\omega}) \\
& + \psi \sum_{t, j, k} |s_{jk, t+1}^{\omega} - s_{jkt}^{\omega}| \\
& + \upsilon \sum_{t, i} P_i \min_{j: x_j=1} (d_{ij} + \vartheta \tau_{ij}^{\omega} + \varrho B_{ij}) \\
& + \omega_1 \sum_{t, j} \max(0, G_j x_j - \bar{G}_j)^2 \\
& + \omega_2 \sum_{t, j} x_j (F_j + H_j + Z_j + A_j) \\
& + \omega_3 \sum_{t, i, k} \left| \frac{\sum_j y_{ijkt}^{\omega}}{D_{ikt}^{\omega} + \varepsilon} - \frac{\sum_j y_{ijk, t-1}^{\omega}}{D_{ik, t-1}^{\omega} + \varepsilon} \right| \Big]
\end{aligned}$$

subject to

$$\begin{aligned}
\sum_{j \in J} y_{ijkt}^\omega + u_{ikt}^\omega &\geq D_{ikt}^\omega, & \sum_{i \in I} \sum_{k \in K} y_{ijkt}^\omega &\leq C_{jt}^\omega x_j, \\
s_{jk,t+1}^\omega &= s_{jkt}^\omega + \text{in}_{jkt}^\omega + \sum_i q_{ijkt}^\omega + r_{jkt}^\omega - \sum_i y_{ijkt}^\omega - \text{loss}_{jkt}^\omega, \\
r_{jkt}^\omega &\leq R_{jkt}^\omega x_j, & y_{ijkt}^\omega &\leq M a_{ij}^\omega x_j, \\
\sum_{j \in J} a_{ij}^\omega x_j &\geq n_{ik}^{\min}, \\
\sum_j y_{ijkt}^\omega &\geq \rho_{ik} D_{ikt}^\omega, & \sum_{i,j,k,t} e_{ijkt}^\omega y_{ijkt}^\omega &\leq E^{\max}, \\
x_j &\in \{0, 1\}, & y_{ijkt}^\omega, q_{ijkt}^\omega, r_{jkt}^\omega, s_{jkt}^\omega, u_{ikt}^\omega &\geq 0.
\end{aligned}$$

The first two terms of the objective measure transport work and the emissions associated with it. The following terms penalize unmet need, unequal service across communities, and excessive travel time to reach a facility. The remaining terms penalize overloaded facilities, dangerous concentration of essential inventories in too few locations, insufficient redundancy across the network, deterioration of goods in storage, ecological disturbance at a site, displacement of existing communities by a facility's footprint, repair activity that exceeds a site's actual repair capability, loss of provenance information as goods move through the network, excessive oscillation in inventory levels from one period to the next, political or jurisdictional distance between a facility and the communities it serves, regional monopoly by a single facility over an essential good, exposure to compound hazards, and abrupt changes in access between consecutive periods. The constraints require that demand be either met or explicitly recorded as unmet, that no facility be loaded beyond its physical capacity, that material be conserved through time net of documented losses, that repair activity remain within a site's actual capability, that emergency redundancy be preserved through a minimum-access-points requirement, that a minimum service floor be met for every community, and that the network remain within an explicit ecological budget.

The Length of the Formula Is Part of the Argument

A distribution site is never simply optimal. It is optimal relative to prior judgments about which distances matter most, whose unmet need is treated as most serious, how much redundancy is prudent against disruption, what ecological disturbance is tolerable at a given site, whether affected communities have standing to contest a proposed facility, and how strongly present transport efficiency should be traded against future resilience. None of the coefficients α through ω_3 in the expression above descend from mathematics. Each one records a political and ethical decision made before optimization begins, and a placement that appears rigorously optimized can nonetheless encode an indefensible set of prior judgments if those coefficients were chosen carelessly, or chosen to favor one class of stakeholder over another under cover of a formula that looks neutral.

This is the same structural point Chapter 1 made about planetary temperature and Chapter 2 made about total biomass, applied now to logistics rather than climate. A single number, whether it is a tar-

get temperature, a biomass total, or a minimized transport cost, can always be produced. Producing it is not the same as having correctly represented the problem, and a placement decision defended only by appeal to a computed minimum has simply relocated the site of judgment into the choice of objective function, where it becomes harder rather than easier to inspect.

What Placement Reveals and What It Should Preserve

A pattern of site-selection decisions, considered across many instances rather than any single case, can disclose more than a facility operator's familiarity with a region. It can reveal which variables are being treated as decisive and which constraints are being honored, whether or not either was ever stated explicitly. A reserve consistently placed where transport routes intersect is implicitly weighting the transport terms of the objective heavily. A reserve consistently placed where repair knowledge already exists in the surrounding population is implicitly weighting repair capability alongside, or above, raw transport efficiency. A reserve consistently placed at a jurisdictional boundary, rather than at the geographic center of its catchment, may be responding to a structural feature of boundaries themselves: a boundary is often where mismatches accumulate, where one region's surplus meets another's shortage, where one production cycle ends as another begins, or where one jurisdiction can decline responsibility for consequences generated elsewhere. A placement network that clusters at such boundaries is not necessarily inefficient. It may be tracking a real feature of where imbalance concentrates that a purely geometric transport-cost minimization would miss.

The question this chapter has been building toward is therefore not simply where material should be placed, but what a given placement preserves. A reserve network that preserves access, reversibility, provenance, repairability, and local agency is not a terminal accumulation of matter awaiting eventual disposal. It is a material memory: a standing reserve of capacity that keeps objects within reach of possible futures rather than converting them, prematurely, into waste. Understood this way, distribution is not the final movement of goods toward a point of consumption after which the material's history ends. It is the continuing placement of matter where some further, presently unspecified use remains admissible, held there against the moment a community's need for it becomes concrete. The remaining chapters of this part extend this standard from a single class of reserve facilities to the transport and routing infrastructure, developed under the name orthodromic infrastructure, that connects them into a planetary network.

Chapter 18

Orthodromic Infrastructure: Routing Between Reserves

Chapter 5 established where distribution reserves belong. It did not establish how material actually moves between them once sited. This chapter takes up that transport layer, under the name this book's own prior development has already given it, orthodromic infrastructure, a routing architecture whose organizing principle is the orthodrome, the great-circle path, the shortest route between two points on a spherical surface, extended from its literal geometric meaning into a governing design principle for planetary-scale material movement: route material along the path that minimizes cumulative cost across the same dimensions Chapter 5's objective function made explicit, not merely the path that minimizes raw distance.

Why Distance Alone Is the Wrong Metric

A naive orthodromic routing scheme would simply compute great-circle distances between every reserve pair and route material along the shortest available path. Chapter 5 already gave the reasons this is inadequate: the transport term in that chapter's objective function was only one of many, alongside emissions, equity, redundancy, spoilage, ecological disturbance, and jurisdictional distance, and a routing layer that optimized distance alone while ignoring the other terms would simply relocate Chapter 5's siting problem one level down, into the routes themselves, rather than resolving it. Orthodromic infrastructure, as this book uses the term, therefore denotes not a literal great-circle router but a routing architecture that treats the great circle as a baseline to be deviated from whenever Chapter 5's other terms justify the deviation, weighted exactly as that chapter's coefficients specify.

This matters concretely for the coupled interventions this book has already evaluated. A route carrying the Caldera Reactor's processed marine carbon output toward a terra preta rain collection basin, the coupling Chapter 17 made explicit, is not well served by the literal shortest sea-to-land path if that path crosses a jurisdiction hostile to the shipment, passes through waters where the Chapter 13 interception fleet's own routes would be disrupted, or requires transshipment through a facility already near the overload threshold Chapter 5's objective penalized. The orthodromic layer's task is to find the route that is shortest *subject to* these constraints, not the route that is shortest *simpliciter*.

Resource Geometry

The infrastructure's flux, Φ , is aggregate corridor throughput capacity, the volume of material a given route can carry per period, constrained by whatever mode, maritime, overland, or a hybrid combining both, the corridor employs. Transport rate, v , is the corridor's actual transit speed, which orthodromic routing treats as a variable to be traded against the other cost terms rather than simply maximized, since the fastest available route is frequently not the route Chapter 5's full objective favors. Storage, S , belongs properly to the reserves this infrastructure connects rather than to the infrastructure itself, though the infrastructure must be sized with reference to reserve storage capacity, a corridor that delivers material faster than a receiving reserve's storage and processing capacity can absorb it simply relocates Chapter 5's overload penalty from the reserve to the corridor's own queuing behavior at the point of arrival.

Admissibility

Physical admissibility, A_P , depends on the specific transport mode a given corridor employs and is not treated uniformly here; maritime routing between coastal reserves is comparatively mature, while routing that must cross difficult overland terrain to reach an inland reserve, or that must interface with a mobile source such as Chapter 13's Sargassum interception fleet or Chapter 14's mobile reactor deployment, requires the corridor itself to be reconfigurable on a timescale shorter than the fixed infrastructure this book has evaluated elsewhere, a requirement more easily stated than engineered.

Institutional admissibility, A_I , compounds the jurisdictional-distance term Chapter 5's objective function already penalized: a corridor, unlike a single reserve, necessarily crosses whatever jurisdictions lie between its endpoints, and every jurisdiction crossed is a further point at which authorization can stall, exactly the multi-jurisdictional exposure Chapters 12 and 13 already identified for mobile oceanic intervention, now generalized to the routing layer connecting every intervention this book has evaluated to every other. Narrative admissibility, A_N , is comparatively favorable for routine cargo movement, a well-understood and broadly accepted category of activity, but inherits whatever narrative burden attaches to the cargo itself: a corridor known to carry Caldera Reactor output may inherit that installation's own narrative profile from Chapter 14, favorable or unfavorable depending on how that profile has developed by the time the corridor is proposed.

Repair

Infrastructure repair for transport corridors follows conventional logistics-network practice, rerouting around a damaged or congested segment while repair proceeds, a form of repair this book's own Chapter 11 already identified as a general advantage of mobility, here extended from the intervention itself to the routes connecting interventions. Institutional repair concerns the jurisdictional stalls described above, and a corridor's design should anticipate, as Chapter 5's own placement analysis argued for reserves, that some jurisdictions along an otherwise favorable route may withhold authorization, requiring either a pre-negotiated framework or a viable alternative routing, ideally both. Model repair concerns the routing algorithm's own weighting of Chapter 5's cost terms against one another, since those weights, as Chapter 5 argued at length, encode political and ethical judgments that must remain open to revision as a population's own priorities shift, rather than being fixed once

at the routing system's initial design and never revisited. Preference repair is comparatively minor for this infrastructure specifically, given its lower narrative visibility relative to the interventions it connects.

Persistence

A corridor network, once established and relied upon, creates a persistence risk closely analogous to Chapter 5's reserve-siting risk: communities and receiving reserves organize their own operations around a given corridor's continued availability, and a corridor's disruption, whether through jurisdictional withdrawal, physical damage, or simple obsolescence as the intervention portfolio it serves evolves, requires the same kind of pre-planned alternative routing this book has argued for throughout, planned before disruption rather than improvised during it.

Collective Choice

The relevant population for any given corridor is the union of populations at both its endpoints and every jurisdiction between them, a population that grows with corridor length and jurisdictional count in a way that makes the civic simulation's role, exposing the corridor's full trajectory and its alternatives to every affected population simultaneously, more rather than less essential as the network scales toward the planetary extent this part's title claims.

Chapter 19

Tracing the Full Material and Energy Chain

The interventions evaluated across Parts III through V, and the reserves and corridors evaluated in this part, have so far been assessed largely as individual nodes: a reactor's own admissibility, a generator's own persistence, a corridor's own repair channels. This chapter asks a different question, one that can only be asked once the routing layer of Chapter 18 exists to connect the nodes: what does the complete chain look like, traced from a given unit of input material or energy through every intervention it passes through, to whatever final, stable, or dissipated state it eventually reaches, and does that complete chain satisfy the reachability-volume objective Chapter 7 established when evaluated as a whole, rather than only at each individual step.

Why Node-Level Admissibility Is Not Sufficient

A chain composed entirely of individually admissible nodes is not guaranteed to be admissible as a chain. This is not a hypothetical caution; it follows directly from the coupling this book has already established between the Caldera Reactor and terra preta rain in Chapter 17. Consider a single unit of Sargassum biomass, intercepted by Chapter 13's fleet, routed via Chapter 18's corridor network to a Caldera Reactor deployment, processed there into biocrude and PHA bioplastic fractions, with some portion of the processed carbon and nutrient output further routed to a terra preta rain collection basin serving a Chapter 16 rainforest generator. Each of these four steps, interception, transport, reactor processing, terrestrial delivery, was separately evaluated for admissibility in its own chapter. But the chain's aggregate energy balance, how much energy the interception fleet consumed, how much the corridor transport consumed, how much the reactor's own operation both consumed and recovered through its turbine, against how much usable energy and stable carbon the chain ultimately delivered to the rainforest system, was not computed at any single step, because no single step's own admissibility evaluation required computing it.

This chapter's method is therefore to trace specific chains explicitly, unit by unit, through every stage, tallying energy consumed and recovered, carbon released and retained, and material transformed at each stage, and to evaluate the resulting aggregate against Chapter 7's reachability-volume standard directly, rather than inferring the chain's acceptability from the acceptability of its individual links.

A Traced Example: Sargassum to Rainforest Carbon

Tracing the chain described above concretely: interception consumes fuel and crew-time proportional to the interception fleet's search-and-recovery distance from base, a quantity that varies substantially with a given bloom's location relative to the fleet's home port, as Chapter 13 already noted. Orthodromic transport from interception point to reactor deployment consumes further fuel, partially offset if the reactor itself has relocated toward the bloom under Chapter 14's mobile-deployment mode, reducing this leg's distance at the cost of the reactor's own repositioning energy, a tradeoff this chapter's tracing method is specifically designed to expose rather than assume favorable in either direction. Reactor processing consumes geothermal-steam energy for the lift-compression cycle, partially recovered through the integrated turbine, and further energy for whichever pyrolysis pathway, hydrous or anhydrous, the incoming material's moisture content routes it toward. Terrestrial delivery of the processed carbon and nutrient fraction consumes further transport energy, offset against the basin's own gravity-fed distribution once material arrives, which this book's Chapter 17 already noted requires no ongoing energy input beyond the initial elevation gained during transport.

The aggregate chain, traced this way, either shows a net energy and carbon balance favorable enough to justify the chain's full complexity against the simpler alternative of leaving the Sargassum to strand and decompose, the maladmitted outcome Chapter 2 diagnosed, or it does not, and this chapter's contribution is not to assert which outcome obtains in general, since the answer depends on the specific distances, fleet efficiency, and reactor operating mode a given real deployment would involve, but to insist that the tracing be performed explicitly, chain by chain, before any specific deployment is authorized, exactly the kind of consequence-bearing exploration Chapter 6 argued a civic simulation must provide in place of a narrative assurance that the chain surely nets out favorably because each of its individual links does.

Generalizing the Tracing Method

The same method applies to every chain this book's interventions can form. A chain connecting Chapter 9's refrigeration installation to whatever grid or fuel source supplies its reactor's own energy input must be traced back to that source's own admissibility, since a refrigeration installation whose energy ultimately derives from a source with its own severe admissibility problems has not resolved the polar margin's viability at the cost of merely relocating an equivalent problem elsewhere, a relocation this book's placement chapter warned against in the material context and that applies with equal force in the energy context. A chain connecting Chapter 12's iceberg-towing operation to the fuel consumed by the towing vessels themselves must be traced against the storm-damage avoided, to determine whether the intervention's own carbon cost is small relative to the benefit it claims, a comparison this chapter's method makes possible and that no single-node evaluation in Chapter 12 was positioned to perform.

This generalization gives the civic simulator specified in Chapter 7 a concrete additional function beyond what that chapter described: not only comparing individual interventions against one another and against their own termination, but tracing the full chains that connect interventions once they are coupled, and evaluating each traced chain's aggregate admissibility as its own distinct object of assessment, since a chain, as this chapter's opening argument showed, is not simply the conjunction of its links' individual admissibility.

Chapter 20

Coupled Intervention Portfolios

The preceding chapter traced individual chains. This closing chapter of Part VI asks the portfolio question directly: given the full set of interventions this book has evaluated, polar albedo maintenance in its two forms, oceanic iceberg routing, kelp and Sargassum processing, the Caldera Reactor, rainforest generators and terra preta rain, and the reserve and corridor infrastructure of Chapters 5 and 18, which combinations of these interventions, pursued together rather than singly, expand the reachable-futures volume Chapter 7 identified as this book's governing objective by more than any of the interventions could expand it alone, and which combinations instead create new admissibility problems at their interfaces that none of the individual interventions faced in isolation.

Complementary Portfolios

The marine-terrestrial coupling of Chapters 14 and 17 is this book's clearest instance of a complementary portfolio, one whose combination this chapter need not re-argue, since Chapter 17 already established it directly. A second complementary pairing, less developed in earlier chapters but visible once the full portfolio is considered together, links Chapter 10's cloud-and-salt hybrid albedo intervention to Chapter 12's iceberg-routing proposal: both address ocean-surface thermal conditions, the former proactively and seasonally at fixed glaciological seams, the latter reactively and briefly ahead of a specific forming storm, and a jurisdiction that has already built the multilateral authorization framework Chapter 12 argued was a precondition for iceberg deployment has, in the course of doing so, likely built much of the institutional infrastructure Chapter 10's own more favorable but still nontrivial admissibility profile would also require, reducing the marginal institutional cost of adopting the second intervention once the first is already authorized.

A third pairing links the reserve and corridor infrastructure of Chapters 5 and 18 to every material-producing intervention this book has evaluated, not as a complementary intervention in its own right but as the substrate every other intervention's material output ultimately depends on for distribution; a portfolio that adopts the Caldera Reactor, or rainforest generators at scale, without first, or at least concurrently, building adequate reserve and corridor capacity is, in effect, adopting an intervention whose own resource-geometry storage term Chapter 7's procedure requires evaluating, and finding that term unfavorable, since processed output with nowhere admissible to go accumulates at the processing site itself, recreating at the reactor or generator's own location the incorporation-capacity failure Chapter 2 first diagnosed at the level of an unmanaged coastline.

Conflicting Portfolios

Not every combination is complementary. Chapter 12’s iceberg-routing proposal depends on a calving supply that Part III’s own framing chapter identified as a symptom of the ice-sheet instability Chapter 9’s and Chapter 10’s polar interventions are meant to arrest; a portfolio that pursues polar intervention aggressively enough to meaningfully slow calving, the success condition those chapters’ own admissibility standards require, correspondingly reduces the iceberg supply Chapter 12’s proposal depends on, a direct tension between two interventions this book has evaluated separately but which a portfolio-level assessment must confront jointly. This tension does not resolve either intervention’s individual case; Chapter 12 was evaluated as a short-horizon, storm-specific mechanism whose ordinary operation does not require the same calving rate its supply depends on to persist indefinitely, but a portfolio planning for both interventions’ long-term coexistence should model this dependency explicitly, inside the civic simulator, rather than treating the two interventions’ resource bases as independent simply because their evaluating chapters were.

A second tension links Chapter 13’s Sargassum interception fleet to Chapter 14’s reactor when the reactor operates in fixed, kelp-farm-moored mode rather than mobile deployment: a fixed reactor cannot process Sargassum wherever it happens to accumulate that season, and a portfolio that adopts the reactor in fixed mode while also committing to Sargassum interception as a standing service, the persistence risk Chapter 13 already flagged, must either accept that some intercepted Sargassum will need long-distance corridor transport to reach the fixed reactor, incurring the energy cost Chapter 19’s tracing method would expose, or accept that the reactor’s fixed siting limits how much of a given season’s bloom it can actually process, a capacity constraint no single chapter’s individual evaluation surfaced because each treated the reactor’s operating mode, fixed or mobile, as a scenario to be selected rather than a portfolio-level commitment interacting with a separately evaluated intervention’s own scaling.

The Portfolio as the Actual Unit of Evaluation

These examples generalize into this chapter’s central claim, and this part’s contribution to the book as a whole: the individual intervention, evaluated in isolation across Parts III through V, is a necessary but not sufficient unit of analysis. The actual decision any population or set of populations faces is not whether to adopt intervention X considered alone, but which portfolio of interventions, given their complementarities and their conflicts, to pursue together, and Chapter 7’s reachability-volume objective, $\text{Vol}_R(A) > 0$, must ultimately be evaluated at the portfolio level, across the full coupled system Chapters 18 and 19 have now given the tools to trace, rather than summed naively across interventions whose individual admissibility, as this chapter’s conflicting-portfolio examples showed, does not guarantee their joint admissibility once their resource dependencies and material chains actually interact.

What Part VI Establishes, and What Remains

Part VI has now supplied the full material-circulation apparatus this book’s remaining argument requires: siting, in Chapter 5; routing, in Chapter 18; chain-tracing, in Chapter 19; and portfolio-level evaluation, in this closing chapter. Every intervention evaluated in Parts III through V can now

be assessed not only on its own terms but as a node in a traceable chain and a member of a jointly evaluated portfolio, exposing complementarities, the marine-terrestrial loop chief among them, and conflicts, the calving-supply tension and the reactor-siting-mode constraint chief among them, that no single intervention's own chapter was positioned to surface. Part VII now turns to the question this portfolio-level apparatus has been assuming a mechanism for but has not yet supplied: given a traced, portfolio-evaluated set of interventions, who is actually authorized to commit to them, on what evidentiary basis, and how can that commitment be terminated, should it need to be, without the termination itself becoming the second crisis Chapter 4 warned every sustained intervention risks creating.

Part VII

Governance, Failure, and Repair

Chapter 21

Record, Recognize, Authorize, Commit, Consequence

Every intervention evaluated in Parts III through VI was assessed as though the question of who may decide to build it had already been settled, so that Chapter 7's five-step procedure could concentrate on resource geometry, admissibility, repair, persistence, and collective choice without also having to specify, case by case, the actual authorization pathway each proposal would need to pass through. That simplification was necessary to keep those parts tractable, but it cannot stand as this book closes. This chapter supplies the authorization pipeline those parts assumed: five distinct stages, record, recognition, authorization, commitment, and consequence, each of which a proposal must pass through, and each of which can fail independently of whether the proposal is admissible in the physical, institutional, or narrative sense Chapter 7 already distinguished.

The Five Stages

A proposal is recorded when its existence, its claimed mechanism, and its supporting evidence are entered into some durable, inspectable form, a published design, a filed patent, a registered feasibility study, available for others to examine independent of whether anyone has yet acted on it. Recording requires nothing beyond documentation; Chapter 9's refrigerator-ice-machine proposal and Chapter 12's iceberg-routing mechanism are both, at the level this book has evaluated them, recorded rather than authorized.

A proposal is recognized when a body with some legitimate claim to evaluate it, a scientific community, a regulatory agency, an affected population exploring it inside the civic simulator Chapter 7 specified, has examined the record and found it credible enough to warrant further attention, without yet granting permission to act. Recognition is where Chapter 7's stratified admissibility test, A_P , A_I , A_N , is actually applied for the first time in a proposal's institutional life; a proposal can be extensively recorded while remaining unrecognized, ignored rather than rejected, which is a distinct and, as the next chapter will argue, easily overlooked failure mode.

A proposal is authorized when the specific body or bodies with jurisdiction over the resources, territory, or population the proposal would affect grant explicit permission for it to proceed, a permission this book has treated throughout as properly belonging with the civic simulation process of Chapter 6 rather than with unilateral declaration. Chapter 9's nuclear-refrigeration proposal was shown to face a comparatively severe authorization barrier precisely because its recognition, on physical

grounds, outpaces its authorization, which remains blocked by the Antarctic Treaty System and Arctic sovereignty disputes; this is the gap Chapter 7 called the misclassification of institutional stall as physical impossibility, now placed at its correct stage in the pipeline.

A proposal is committed when authorization converts into actual resource allocation, construction, deployment, or operation, the point at which Chapter 4's reachability and viability standards begin to apply to the real world rather than to a simulated exploration of it, and the point past which Chapter 6's distinction between free exploration and costly commitment takes hold.

A proposal reaches consequence when its committed operation, and eventually its termination, produces the actual outcomes, favorable, unfavorable, or mixed, that the rest of this book's apparatus, resource geometry, repair, persistence, has been built to evaluate and anticipate. Consequence is not a final stage in the sense of closing the pipeline; it feeds back into the record, since a proposal's actual observed consequences become part of what future recognition decisions, for that proposal's continuation or for a related proposal's initial evaluation, must take into account.

Why the Stages Must Be Kept Distinct

The temptation this chapter exists to resist is collapsing these five stages into a single question, has this proposal been approved, since that collapse obscures exactly where a given proposal's actual difficulty lies. A proposal well recorded and well recognized but stalled at authorization, Chapter 9's case, has a different problem, and needs a different remedy, institutional repair in Chapter 7's sense, than a proposal that has never been adequately recorded or recognized in the first place, whose problem is visibility rather than permission. Conflating the two produces exactly the misdiagnosis Chapter 7 warned against: treating an institutional or evidentiary gap as though it were a physical one, or treating a genuinely unresolved physical question, Chapter 10's contested cloud-brightening admissibility, as though it were merely awaiting bureaucratic sign-off.

Chapter 22

Evidence Before Recognition: What Survives Is Not What Was Best

The pipeline Chapter 21 specified has a structural property this book has so far used without examining directly: at every stage, the population of proposals shrinks, and it shrinks for reasons that are only partly related to the proposals' actual admissibility in Chapter 7's sense. This chapter examines that shrinkage directly, because failing to examine it risks a specific and consequential error: mistaking the proposals that happened to survive the pipeline for the proposals that were actually best suited to expanding the reachable-futures volume Chapter 7 identified as this book's governing objective.

The Population of Unbuilt Interventions

Let G denote the full population of geoengineering proposals that have ever been conceived, at whatever level of seriousness, for addressing the problems this book's Part I diagnosed. Most of G never reaches the recording stage Chapter 21 specified: an idea discussed once at a conference, sketched on a napkin, mentioned in a grant application that was never funded, and abandoned. Call the recorded subset $G_1 \subset G$. Of G_1 , a smaller subset G_2 is recognized, examined seriously enough by some credentialed body to generate a documented assessment of its admissibility. Of G_2 , a smaller subset G_3 is authorized, granted permission by whatever jurisdiction its deployment would require. Of G_3 , a smaller subset G_4 is actually committed, built or deployed at least at pilot scale. And of G_4 , this book itself has only had occasion to evaluate a small subset, G_5 , the handful of proposals, refrigeration-based ice maintenance, cloud-and-salt hybrid albedo, iceberg routing, kelp and Sargassum processing, the Caldera Reactor, rainforest generators, terra preta rain, whose visibility and documentation were sufficient for Parts III through V to treat them as this book's worked examples.

The population this book has actually reasoned about, in other words, is G_5 , a set selected by five successive filters, only some of which have anything to do with the continuation-geometric admissibility standard Chapter 4 established. A proposal can fail to reach G_1 because its originator lacked the institutional access to document it credibly, not because the mechanism it describes is physically unsound. A proposal can fail to reach G_2 because it was proposed at an inopportune moment, competing for attention against a more narratively compelling rival exactly of the kind Chapter 6 described, not because a careful admissibility review would have rejected it. A proposal can fail to reach G_3 because the jurisdiction with authority over it changed political composition before authorization could be secured, an accident of timing rather than a judgment about the proposal's merit.

A proposal can fail to reach G_4 because funding collapsed for reasons entirely exogenous to the proposal itself, a war, a recession, a change in a funding agency's institutional priorities. And a proposal can fail to reach G_5 , this book's own attention, simply because its documentation did not happen to surface during this project's own research, a filter with no claim whatsoever to tracking admissibility.

The Error of Reasoning Backward from Survivors

The danger this book must guard against, having now made this filtering structure explicit, is reasoning backward from G_5 as though its members' survival to this stage were itself evidence of their superior admissibility. This is a natural inference and a mistaken one. A proposal that reached pilot deployment did so because it survived recording, recognition, authorization, and commitment, and each of those stages, as the preceding section showed, filters on funding continuity, institutional access, narrative timing, and simple visibility, in addition to whatever it filters on admissibility. A population of surviving proposals is therefore not a sample of the best proposals in G ; it is a sample of the proposals that were simultaneously admissible enough, and fortunate enough in funding, timing, and visibility, to survive five successive, only partially correlated filters.

This has a direct and uncomfortable consequence for how this book's own worked examples should be read. Chapter 14's Caldera Reactor and Chapter 16's rainforest generators are evaluated here in detail not because a rigorous prior screening of all conceivable marine-processing or terrestrial-support architectures established them as the two best available designs, but because they were the designs this project's own development had already produced and documented in sufficient depth to serve as worked examples. The vast invisible population of unbuilt, undocumented, or simply unencountered alternative designs, some of which may dominate these two on Chapter 7's reachability-volume objective, is not represented in this book at all, and this book's own selection of examples is therefore itself an instance of exactly the survivorship the population G_5 describes.

What This Means for Evidence Standards at Recognition

The practical consequence for Chapter 21's authorization pipeline is that the recognition stage specifically, the point at which a body first evaluates a proposal's admissibility, must ask a question beyond the immediate one Chapter 7's stratified test already poses. It must ask not only whether this specific proposal satisfies A_P , A_I , and A_N , but whether the proposal's presence at the recognition stage at all reflects a filtering process a population evaluating it should trust. A proposal that reached recognition because a well-resourced institution with existing access to regulatory bodies pushed it forward carries a different evidentiary weight than a proposal that reached recognition on the strength of its documented mechanism alone, even if both proposals, once actually evaluated against Chapter 7's admissibility test, score identically. The civic simulator specified in Chapter 7 should therefore be understood as having a further function beyond the ones that chapter already described: not only exploring a given proposal's trajectory, but exposing, where the information exists, how that proposal came to be under consideration at all, so that a population is not systematically evaluating only the recorded, recognized, and visible fraction of G while remaining unaware of what filtering produced that fraction.

This does not mean every proposal in the vast unrecorded remainder of G deserves equal considera-

tion; most of G almost certainly fails Chapter 7's admissibility test for perfectly good reasons, just as most creative works that never reach an audience are not secretly superior to the ones that do. It means only that the fraction of G a population actually gets to evaluate has already been shaped by filters uncorrelated with admissibility, and that a rigorous recognition process should ask about that shaping explicitly rather than treating a proposal's mere presence at the recognition stage as itself a credential.

Procedural Leverage and the Interventions Already Evaluated

A related distinction is worth drawing out explicitly, since it bears on how this book's own worked examples should be weighed against whatever unrecorded alternatives exist in G . Some of the labor that produced G_5 's members was, in an important sense, leveraged rather than merely accumulated: a design choice, the kelp-baseline allocation of Chapter 13, the dual hydrous-anhydrous pyrolysis split of Chapter 14, that resolves a whole class of subsequent problems at once is not doing the same kind of work as a design choice that merely adds one more increment of engineering detail. This book has favored, throughout, interventions whose design choices operate this way, reducing rather than merely relocating a downstream problem, exactly the standard Chapter 19's chain-tracing method used to distinguish a genuinely favorable material and energy chain from one that only looked favorable because its costs had been pushed to a different link. A population evaluating candidate interventions at the recognition stage should ask, of any proposal, whether its central design choices are of this leveraged kind or whether they are merely additive, since a proposal built on leveraged design choices is more likely to survive the kind of unanticipated perturbation Chapter 7's six-posture stress test was built to expose, having already demonstrated that it resolves classes of problems rather than individual instances of them.

What Part VII Still Owes

This chapter has established that the authorization pipeline is not evidence-neutral: what survives to recognition, and what survives further to commitment, reflects funding, timing, and visibility as much as it reflects the continuation-geometric merit this book's apparatus is built to assess. The remaining chapters of this part take up the two questions this filtering structure makes more urgent rather than less: how a committed intervention's eventual termination can be planned for without becoming, as Chapter 4 warned, a second crisis in its own right, and what a society actually learns, and should institutionally retain, from the interventions in G_5 that were tried and that failed, given that this chapter has just shown the population of tried-and-failed interventions is itself a biased sample of a much larger population of untried ones.

Chapter 23

Termination Without Termination Shock

Chapter 4 introduced termination shock as the paradigmatic risk this book's continuation-geometric standard is built to anticipate: a sustained intervention's surrounding systems reorganize around its permanent presence, so that its withdrawal, whether planned, forced, or accidental, produces a second crisis rather than a restoration of the state that preceded the intervention. Every intervention evaluated in Parts III through V was assessed, in its persistence subsection, for this risk specifically. This chapter asks what an adequate termination plan actually requires, now that the full authorization pipeline of Chapter 21 and the selection-effect caution of Chapter 22 are both in view, and argues that termination planning is not a late-stage addendum to an intervention's design but a condition that should be evaluated at the recognition stage, before authorization, before commitment, alongside every other admissibility criterion Chapter 7 specified.

Why Termination Planning Cannot Wait Until Termination

A termination plan drafted after an intervention has already reorganized its surrounding systems around its presence is not a termination plan in any meaningful sense; it is a crisis-response document written under exactly the conditions, urgency, incomplete information, institutional pressure, that produce the worst outcomes. Chapter 9's nuclear-refrigeration installation illustrates the point sharply: a decommissioning plan for a reactor that has operated for decades, adjacent to a coastal or ecological system that has adapted to the margin it maintained, cannot be improvised at the point funding collapses or political support withdraws. It must be specified, in enough detail to be evaluated against Chapter 4's recoverability standard, before the installation is ever authorized, so that the population granting authorization is granting it with full knowledge of what withdrawal would require and how long it would take.

This requirement changes what counts as a complete proposal at the recognition stage. A proposal that specifies an intervention's operating mechanism, expected benefit, and estimated cost, but leaves its termination trajectory unspecified, is not a complete admissibility case in this book's sense; it has addressed Chapter 7's resource-geometry and much of its admissibility analysis, but has left the persistence analysis, one of Chapter 7's five explicit steps, incomplete. This book's own worked examples were held to exactly this standard throughout Parts III through V, and this chapter now states as a general principle what those chapters demonstrated case by case: a persistence analysis without a specified termination trajectory is not a persistence analysis.

Graded Termination and the Autonomy Standard

Chapter 16's rainforest generator supplied this book's clearest example of what an adequately specified termination trajectory can look like when an intervention's dependency is graded rather than binary. Because that chapter tracked the fraction of supported biomass that had transitioned to ecological autonomy, its termination trajectory was not a single cliff-edge event but a monitored variable, permitting a population to know, at any point in the intervention's operation, how severe a termination shock would currently be if withdrawal became necessary, and to weigh that severity against how slowly or quickly the autonomy fraction was rising. This graded structure generalizes: wherever an intervention's dependency can be measured continuously rather than assumed absent until a single catastrophic failure, its termination planning should track that measure explicitly, converting an otherwise binary and unpredictable shock into a monitored and, to some degree, manageable one.

Not every intervention this book has evaluated admits this graded structure equally well. Chapter 9's installation either operates or does not; there is no intermediate fraction of the margin that has become independently self-sustaining the way Chapter 16's forest biomass could. For interventions of this binary kind, termination planning must instead specify a wind-down sequence, a staged reduction in operating capacity paired with staged monitoring of the margin's own response, rather than an autonomy fraction, but the underlying requirement, that the trajectory be specified and monitorable rather than assumed, is the same.

Termination and the Authorization Pipeline

Placing termination planning at the recognition stage, rather than treating it as a condition of authorization alone, has a further consequence for Chapter 21's pipeline: a proposal's termination trajectory should itself be recorded, alongside its operating mechanism, and should be revisable through Chapter 7's model-repair channel as operating data accumulates, exactly as the operating mechanism's own performance is. A termination plan drafted once, at initial recognition, and never revisited as the intervention's actual surrounding dependencies develop in ways the original plan did not anticipate, is vulnerable to exactly the kind of stale, unrevised model Chapter 7 warned against throughout. Chapter 13's kelp-and-Sargassum system illustrates why this matters concretely: a termination plan for the Sargassum interception service, drafted before any bloom season had actually been serviced, could not have anticipated the specific coastal communities' actual adaptation to the service's presence, which Chapter 13's persistence analysis flagged as a risk precisely because it would only become visible after the service had operated for some years. An adequate termination plan for that service must therefore be a living document, updated against observed community dependency, not a fixed specification frozen at the point of initial authorization.

Multilateral Termination for Multilateral Interventions

Chapter 12's iceberg-routing proposal and Chapter 18's corridor infrastructure both raised, in their own admissibility analyses, the problem of authorization across multiple jurisdictions. Termination inherits this multiplicity directly: an intervention authorized by several jurisdictions jointly cannot be terminated by any single one of them unilaterally without risking exactly the kind of unilateral-action problem this book's Introduction used Stephenson's Schmidt to dramatize, now applied to

withdrawal rather than initiation. A multilaterally authorized intervention therefore requires a multilaterally specified termination protocol, agreed upon at the same time as the initial authorization, specifying not only the technical wind-down sequence but which body or bodies must jointly approve its initiation, so that withdrawal cannot become a unilateral act any more than the original commitment was meant to be.

What Adequate Termination Planning Requires

Drawing the preceding sections together, this book proposes that a termination plan, to be considered adequate at the recognition stage of Chapter 21's pipeline, must specify four things: a monitorable measure of the surrounding systems' dependency on the intervention's continued operation, graded where the underlying dependency is graded and staged where it is not; a wind-down sequence whose steps are evaluated against Chapter 4's recoverability standard at each stage, not only at the sequence's completion; a revision protocol permitting the plan itself to be updated as observed dependency diverges from the plan's original assumptions, following Chapter 7's model-repair channel; and, for any intervention authorized by more than one jurisdiction, a joint termination-initiation protocol precluding unilateral withdrawal in the same way this book has argued throughout against unilateral commitment. An intervention whose recognition-stage proposal cannot supply these four elements has not completed Chapter 7's persistence analysis, regardless of how favorably its operating mechanism otherwise scores.

Chapter 24

Repair as Planetary Knowledge

The chapters of this part have so far treated repair, in Chapter 7's four-channel sense, infrastructure, institutional, model, and preference, as a capacity each intervention must possess. This closing chapter asks a different question: what does a population, or a civilization, actually learn from an intervention's failure and subsequent repair, and how should that learning be retained so that it improves the evaluation of the next proposal entering Chapter 21's pipeline, rather than being relearned, at cost, each time a structurally similar failure recurs.

Failure as the Book's Own Missing Category

It is worth noting directly what this book has not done. Every intervention evaluated in Parts III through V was assessed prospectively, before construction, using Chapter 7's five-step procedure and six-posture stress test. None of them has actually failed, in the sense of having been built, operated, and then required the repair machinery this book's apparatus specifies, because none of them, at the level of detail this book has evaluated, has yet been built. This is not a defect in the argument; it follows directly from Chapter 22's own point that this book's worked examples are drawn from G_5 , the small population of proposals visible and documented enough to serve as examples, most of which have not yet progressed past Chapter 21's authorization stage, let alone reached the consequence stage where failure and repair would actually occur. But it means this chapter's treatment of repair as planetary knowledge must necessarily be more schematic than the preceding chapters' treatment of prospective admissibility, since it describes a capacity this book's own examples have not yet had occasion to exercise.

What a Repair Record Must Preserve

Chapter 7 specified that model repair should maintain a public, auditable provenance record of each revision, distinguishing a model's history of correction from a record that quietly discards its own past errors. This chapter generalizes that requirement from individual models to the full authorization pipeline: a failure at any stage, an intervention that reached commitment and then required infrastructure, institutional, model, or preference repair, should generate a record specifying which admissibility dimension actually failed, A_P , A_I , or A_N , which repair channel was actually required to correct it, and how long the correction took relative to the original persistence analysis's estimate. This record is what Chapter 22's selection-effect caution requires to be usable at all: a future recognition-stage evaluation of a structurally similar proposal should be able to consult not only

whether prior similar interventions survived to commitment, the biased signal Chapter 22 warned against, but what specifically went wrong with those that failed after commitment, information the pure survivorship record does not carry.

This distinction matters because a failure record and a survival record answer different questions. A survival record answers which interventions made it through Chapter 21's pipeline; a failure record answers, of those that made it through and were then operated, which specific admissibility dimension proved to have been misjudged, and by how much. Chapter 22 already showed why the survival record alone is an unreliable guide to a proposal's actual merit. The failure record is this book's proposed remedy: a durable institutional habit of documenting not only successes but the specific, dimension-attributed nature of failures, retained and consulted at every subsequent recognition-stage evaluation.

Repair Knowledge Across the Portfolio

Chapter 20's portfolio-level analysis already showed that individual interventions' admissibility does not guarantee their combined admissibility once coupled; the calving-supply tension between Chapter 12's iceberg proposal and Chapter 9's or Chapter 10's polar interventions, and the siting-mode conflict between Chapter 13's interception fleet and Chapter 14's reactor, were both surfaced only by tracing the portfolio as a whole. Repair knowledge inherits this same requirement. A failure that occurs because two jointly committed interventions interact in a way neither's individual recognition-stage evaluation anticipated is not adequately recorded as a failure of either intervention considered alone; it must be recorded as a portfolio-level failure, attributed to the specific coupling that produced it, so that a future population considering a similar coupling, even between two entirely different interventions that happen to share the same structural dependency, calving supply and active removal, or fixed processing capacity and mobile feedstock, can consult the prior portfolio-level failure rather than rediscovering the same coupling problem independently.

Repair Knowledge and the Civic Simulator

Chapter 7 specified that the civic simulator's governance must be endogenous, its own rules for modification and revision built into the simulated institutional space rather than imposed from outside. This chapter's repair record belongs inside that endogenous structure directly: the simulator's own model of a given class of intervention, ice-margin maintenance, biomass processing, terrestrial support architecture, should be revised, with public provenance, each time a real deployment's failure record becomes available, so that the simulator a future population uses to explore a new proposal in the same class is already informed by every prior deployment's actual, dimension-attributed failure history, rather than only by the prospective admissibility analysis this book's own chapters have been limited to providing. This is, in the end, what separates a civilization capable of learning, in Chapter 7's own terms, from a hypothesis capable of learning from its own consequences, from a civilization that merely accumulates a longer list of things it has tried.

Closing Part VII

Parts III through VI evaluated interventions prospectively. Chapter 21 supplied the pipeline those interventions must pass through to move from proposal to reality. Chapter 22 showed that the population of proposals a society actually gets to evaluate at each pipeline stage is a biased sample, shaped by funding, timing, and institutional access as much as by genuine admissibility, and that this book's own worked examples are themselves an instance of that bias. Chapter 23 argued that termination planning belongs at the recognition stage rather than as an afterthought to authorization, and specified what an adequate termination plan requires. This closing chapter has argued that failure, once it occurs, must be recorded in a form that preserves more than the fact of failure, the specific admissibility dimension and repair channel involved, so that the pipeline's own operation improves with each cycle rather than repeating the same misjudgments under new names. The Conclusion now closes the book by returning to the question the Introduction opened with, what a defensible geoengineering practice actually looks like, once temperature and biomass have been rejected as sufficient objectives, once simulation has replaced narrative as the substrate for evaluating interventions, once specific polar, oceanic, terrestrial, and circulatory interventions have been evaluated in detail, and once the authorization, evidence, termination, and repair apparatus of this closing part is understood as inseparable from the interventions it governs.

Conclusion: Reversible Supports for Plural Continuations

The Introduction opened with a single image this book has spent its remaining length replacing: a machine, or a coordinated set of machines, acting on a single planetary variable to move it toward a preferred value. Stephenson's Schmidt, forcing every affected government and population to respond to a fact his installation had already made irreversible, dramatized what that image costs when it is actually built. This book has not argued that geoengineering should therefore be abandoned. It has argued that the image itself, one apparatus, one variable, one planet moved toward one target, was never an adequate description of what a defensible intervention requires, and it has spent seven parts building the alternative image out in full.

What the Argument Established

Part I dismantled the thermostat model on its own terms. Chapter 1 showed that no optimal global temperature exists to be targeted, because the systems a temperature might be optimized for, biomass, biodiversity, agriculture, coastal stability, ocean productivity, respond to temperature differently and do not share a maximum. Chapter 2 showed that the most obvious fallback, maximizing total biomass, fails for the same structural reason: biomass existing in the wrong place, concentration, or state can represent damage rather than health regardless of its quantity, as the paired krill and Sargassum cases demonstrated with unusual clarity. Chapter 3 reframed climate risk itself, from a story about endangered coastlines to the more general and more accurate story of systems committed to climatic regularities that a changing climate is revising faster than those commitments can be renegotiated. Chapter 4 then supplied the standard the rest of the book would use in place of a single scalar: reachability, viability, recoverability, and path dependence, culminating in termination shock as the paradigmatic risk any sustained intervention creates for itself.

Part II supplied the method by which that standard could actually be applied before, rather than after, an intervention's consequences accumulate. Chapter 6 showed why narrative coordination, discourse, persuasion, symbolic alignment, is structurally unsuited to evaluating high-dimensional, path-dependent interventions, because narrative form survives only by discarding the dimensionality that determines what will actually happen. Chapter 7 supplied the architecture that could do what narrative could not: a civic simulator treating civilizations and their interventions as executable hypotheses, evaluated through stratified admissibility, a five-step explanatory procedure, and a six-posture stress test designed to expose which governing logic a proposal actually embodies rather than merely the case its proponents prefer to tell.

Parts III through VI then applied that method to four families of intervention in turn. Polar albedo maintenance, in its nuclear-refrigeration and cloud-and-salt-hybrid forms, showed that concentrated

leverage and distributed reversibility trade against each other, and that neither proposal dominates the other once repair burden and institutional exposure are weighed alongside operating benefit. Oceanic continuation engineering showed that mobility itself is a resource, purchased at the cost of continuously renegotiated rather than single-site admissibility, across iceberg routing's unresolved kinship with unilateral intervention, the kelp-and-Sargassum system's asymmetric institutional footprint, and the Caldera Reactor's still-undemonstrated engineering synthesis. Constructed terrestrial productivity showed that fixed, durable substrate trades mobility for a graded persistence risk tied to a support architecture's own success in making itself unnecessary, the ecological-autonomy standard Chapter 16 tracked continuously rather than claimed once. And planetary material circulation showed that admissibility evaluated node by node does not guarantee admissibility evaluated as a chain or a portfolio, exposing, once the whole system was traced, two genuine tensions, the calving-supply dependency between polar and iceberg interventions, and the siting-mode conflict between fixed processing and mobile feedstock, that no single intervention's own chapter could have surfaced alone.

Part VII closed the loop those four parts left open. Chapter 21 supplied the authorization pipeline every proposal must actually pass through, recording, recognition, authorization, commitment, consequence, distinguishing failures of institutional permission from failures of physical soundness. Chapter 22 showed that the population of proposals visible at any pipeline stage is a biased survivor sample, shaped by funding, timing, and institutional access as much as by genuine admissibility, and turned that caution on this book's own worked examples rather than exempting them. Chapter 23 argued that termination planning belongs at recognition, not as a crisis document written under the pressure of an intervention's actual withdrawal. Chapter 24 closed by distinguishing a survival record from a failure record, arguing that only the latter, dimension-attributed and portfolio-aware, gives a civilization anything worth calling institutional memory rather than a longer list of things it has tried.

What Success Looks Like

Given all of this, what would it mean for a geoengineering practice built on this book's apparatus to succeed. Not, as Chapter 1 rejected from the outset, the achievement of some single preferred planetary temperature or maximized biomass total. Success, on the standard this book has defended throughout, is the construction of reversible supports for plural continuations: interventions whose operation expands, or at minimum does not contract, the volume of reachable, viable, and recoverable futures available to the populations and ecosystems they touch, and whose own eventual termination has been planned with the same seriousness as their operation, so that withdrawal never becomes the second crisis Chapter 4 identified as every sustained intervention's central risk.

This standard is deliberately not a single number. It cannot be reduced to one, for the same reason Chapter 1 showed no single temperature and Chapter 2 showed no single biomass total could serve as an adequate objective. A successful portfolio of interventions, evaluated by this standard, is one that a population, having explored its trajectories inside the civic simulator Chapter 7 specified, chooses to continue inhabiting, developing, and defending once its consequences are no longer hypothetical but lived, precisely the preference-as-functional-over-trajectories account Chapter 6 argued was the only adequate replacement for preference declared in advance and unrevised by what navigation actually reveals.

What This Book Has Not Done

Honesty about the limits of this project requires stating plainly what it has not accomplished. This book has evaluated a specific, finite set of proposals, drawn overwhelmingly from one research program's own documented work, and Chapter 22's own argument requires acknowledging that this set is not a representative or exhaustive sample of the much larger population of conceivable interventions, most of which remain unrecorded, unrecognized, or simply unencountered by this project. None of the interventions evaluated here has actually been built, operated, and repaired at the scale this book's chapters describe; the resource-geometry, admissibility, and stress-test analyses throughout are prospective, and Chapter 24 was explicit that a genuine failure record, the kind of institutional knowledge this book has argued is indispensable, does not yet exist for any of them. This book is, in its own vocabulary, a recognition-stage document, not a consequence-stage one, and it should be read, and used, accordingly.

The Standard That Remains

What this book has offered, in place of a finished blueprint, is a standard, and a method for applying it, that does not depend on any single intervention it has evaluated turning out to be correct. The ice-albedo control-surface logic of Part III, the mobility-versus-admissibility tradeoff of Part IV, the graded-autonomy persistence standard of Part V, the chain-and-portfolio tracing of Part VI, and the pipeline-and-repair apparatus of Part VII would all remain useful even if every specific proposal this book named, nuclear refrigeration, cloud-and-salt hybrid albedo, iceberg routing, the Caldera Reactor, rainforest generators, terra preta rain, orthodromic infrastructure, were eventually superseded by better designs drawn from the vast, currently invisible population Chapter 22 insisted this book acknowledge. The standard is not which future Earth should have. It is which futures remain reachable after a given intervention's choices have been made, evaluated honestly, revised continuously, and never mistaken for a question a single number could answer.