



# Dangerous-Sounding Speech

## Semantic Proximity and the Inversion of Safety



# The Working Assumption vs. The Reality



**Takeaway:** This is **not random classifier noise**. A noisy classifier errs in both directions. This is a **signed error** that actively **inverts the safety** ranking.

Euphemistic/Coded threats: Veiled threats, coordinated instructions, vague wellness pitches.

Y-Axis  
Causal Contribution  
(Protective/Mitigating to Harm-Advancing)

The Attacker's Escape  
(Safe Zone)

The Protector's Trap  
(Flagged Zone)

Explicit/Clinical utterances:  
Overdose checklists,  
anatomical consent,  
survivor threat logs.

X-Axis  
Semantic Proximity (Lexically Calm to Lexically Alarming)

## Layer 1: Lexical Content


Cheap, scalable, computable.  
(e.g., drug names, slurs, exploits).

## Layer 2: Speech Act

Probabilistically inferable.  
(e.g., warning, quoting, mocking).


## Layer 3: Operational Role

Distributed and non-local.  
(e.g., giving means, withholding opportunity).



**Systems under review  
pressure collapse all  
layers** into the one they  
can cheaply compute.

**Operational role is  
entirely missed** because  
it sits outside the isolated  
item being classified.




# The Governing Principle

- A rule keyed to operational role depends on external, unobservable facts. It cannot be easily defended on appeal. 
- A rule keyed to lexical content can be written down, applied by thousands of reviewers, and defended. 

**Takeaway:** The selection pressure against explicit protection is a structural output of optimizing for legal and operational defensibility. It will recur under any moderation regime that keeps this target.

|                     | Domain            | THE PROTECTOR'S BURDEN<br>(Explicit speech flagged) | The Attacker's Advantage<br>(Coded speech passes) |
|---------------------|-------------------|-----------------------------------------------------|---------------------------------------------------|
| PUBLIC HEALTH       | Harm Reduction    | Naloxone/fentanyl instructions                      | Vague anti-drug slogans                           |
|                     | Sexual Health Ed  | Anatomical vocabulary                               | Coercive, exploitative messaging                  |
|                     | Medical Diagnosis | Naming frightening possibilities                    | Ineffective remedy pitches                        |
| SECURITY & VIOLENCE | Cybersecurity     | Defensive proof-of-concept exploits                 | Fake researcher profiles                          |
|                     | Domestic Abuse    | Itemized survivor accounts                          | Veiled threats with location sharing              |
|                     | Extremism         | Historical documentation of symbols                 | Coded nostalgic imagery                           |

## THE EVIDENTIARY BLACKHOLE




**Human Rights Watch:**  
619 of 5,396 pieces (11%)  
of social media content  
used to support abuse  
and war crime  
allegations were removed  
because naming violence  
explicitly triggers filters.

## THE RESEARCHER'S PENALTY



**Extremism Documentation:**  
Investigative journalist  
Alexa O'Brien was banned  
from YouTube and Gmail  
for hosting al-Qaeda  
videos used as evidence  
in the Chelsea Manning  
trial.

## THE EDUCATIONAL ERASURE



**Meta Oversight Board:**  
Sexual health and gender  
identity education  
removed under Adult  
Nudity policies  
because the rules  
blindly capture precise  
anatomical vocabulary.

*Free adaptation.  
Strategy already in use.*

Platform Warning:  
"Rephrase. Speak  
less explicitly."

Moves to coded language → Evades filter → Goal Advanced

Forced into vagueness → Loses clinical  
precision → Goal Disabled

*Corrosive. Protection requires precision.*

False Positives vs. Adversarial Asymmetry. Asking both parties to phrase things "more safely" is an intervention that disarms the protector and formalizes the attacker's evasion strategy.

**Moderation** selects for actors who know how to make dangerous things sound ordinary. It acts as a filter correlated with rhetorical class posh position, penalizing those who lack the privilege or composure to use institutional indirection.

**Moderation** selects for actors who know how to make dangerous things sound ordinary. It acts as a filter correlated with rhetorical class posh position, penalizing those who lack the privilege or composure to use institutional indirection.

# A content moderation regime built on semantic proximity is not simply biased; it is structurally regressive.

- It taxes the vulnerable who require explicit speech to survive or document harm.
- It actively trains malicious actors to become the exact population the system is worst at reading.
- The error compounds round after round, for as long as defensibility is optimized over accuracy.